
WorldEpisode: Auditing Silent State Drift in Robot-Learning Datasets

Anonymous Author(s)

Affiliation

Address

email

Abstract

Robot-learning datasets can be locally valid and behaviorally wrong: every file loads, but a command is replayed at the moment it was issued rather than when the motor executed it, the same reconstructed scene appears in both training and testing, or an asset’s collision role no longer matches the state under evaluation. We call this *silent state drift*. This paper presents *WorldEpisode*, an executable state-integrity contract for robot-learning datasets: five machine-checkable invariants per episode (persistent identity, frame and clock mappings, representation roles, temporal events, provenance) plus the two corpus-level obligations that carry the headline results, lineage-disjoint evaluation splits and loss-explicit conversion.

On a public SO-101 LeRobot release, a random episode split overlaps every tested task–scene proxy lineage. Across 5 optimization seeds, mean episode nRMSE is 0.183 [0.171, 0.196]; holding out 2 task–scene groups reduces measured proxy-lineage overlap to 0.000 and increases nRMSE to 0.381 [0.364, 0.399], a 0.197 increase [0.177, 0.219]. Because task identity is part of the lineage key, this result does not isolate scene leakage from task shift. A 3-frame lag frozen on 320 trajectories reduces held-out pooled action/state telemetry RMSE from 4.845 to 1.934 source position units on 80 trajectories; the mean paired episode improvement is 3.087 [2.890, 3.286]. This is a telemetry proxy spanning 1 controller configuration and without motor-effective timestamps. A complete-shard round-trip audit over 271 public LeRobot episodes from 3 datasets preserves 43,601 paired state/action samples with maximum numerical error 0.0 while reporting source-absent semantics as conversion loss; the validator detects all 14 injected fault classes. We claim neither simulator-independent physics nor benchmark inflation without faithful policy reruns.

1 Introduction

Robot-learning data can be structurally valid yet behaviorally wrong. A tensor can have the expected shape but an undeclared unit; a command timestamp can be mistaken for motor-effective time; an asset can retain its appearance while losing its collision role; or an episode-level split can put related world captures in both training and evaluation. Each artifact still loads, so ordinary schema and container checks miss the failure. We call this class *silent state drift*: the stored observation, action, asset, or split no longer denotes the physical state assumed by training, replay, or evaluation.

Robot learning is a demanding test case because it combines asynchronous control, multi-camera recordings, reconstructed scenes, simulator exports, and policy evaluation. LeRobot integrates collection, datasets, policies, and deployment [4]; LeRobotDataset v3 stores state, action, timestamps, video, and episode indexes [7]. Rerun and NCore cover multimodal logs and pose conventions, while OpenUSD and glTF describe worlds and assets [2, 15, 19]. These containers are complementary, but

Table 1: The contract at a glance: each question a valid episode must answer, the graph or obligation that answers it, the executable requirement IDs, and the evidence case that tests it.

Question	Answered by	Requirements	Tested in
Which state revision am I in?	world revision + ordered deltas	WORLD.001, TRACE.001	Case 4
What entities does it contain?	identity graph	ENTITY.001	Case 4
Where and when is it valid?	frame + clock graphs	FRAME.001-2, TIME.001-2, ACTION.001-2	Cases 2, 4
What is each asset safe for?	representation-role graph	REP.001, ASSET.001-2	Cases 3, 4
What actually happened?	temporal-event graph	TRACE.001	Case 4
Where did this come from?	provenance graph	PROV.001, QUALITY.001	Case 4
What was lost in conversion?	obligation: loss-explicit conversion	CONVERT.001	Cases 5, 6
Can the evaluation split leak?	obligation: lineage-safe splits	SPLIT.001	Cases 1, 7
Can this scale past one folder?	dataset manifest + indexes	DATASET.001	App. F

37 none by itself proves that the same entities, frames, clocks, action semantics, asset roles, and source
38 lineages survive conversion.

39 **Related work.** Documenting datasets is not new. Datasheets for Datasets [5] and model cards [14]
40 standardize prose descriptions of provenance and intended use; MLCommons Croissant [1] adds a
41 machine-readable metadata layer for ML-ready datasets; and Open X-Embodiment [17] aggregates
42 robot data under a shared schema. Evaluation leakage is a well-documented driver of the reproducibil-
43 ity crisis in ML-based science [8]. WorldEpisode differs by making cross-container assumptions
44 executable. It binds each episode to an immutable world revision and closed identity, frame/clock,
45 representation-role, event, and provenance graphs; a fail-closed validator reports stable requirement
46 IDs, and each conversion emits machine-readable loss. It is a sidecar contract over existing containers,
47 not a replacement format or dataset card.

48 Contributions.

- 49 1. An executable state-integrity contract with five coupled graphs, JSON Schema, semantic validator,
50 conformance fixtures, and fail-closed LeRobot/Rerun preflight.
 - 51 2. Loss-explicit LeRobot conversion, including exact round trips over 271 public episodes from 3
52 datasets and a versioned binding-retention audit.
 - 53 3. Empirical audits of a task–scene-proxy split and action timing, with episode-level uncertainty
54 intervals and explicit boundaries for the confounded split, single-controller telemetry, same-trace
55 simulator replay, and controlled proxy experiments.
- 56 Section 2 states the contract, Section 3 makes it executable, Section 4 reports the audits, and Section 6
57 states their limits.

58 2 The Contract: Five Graphs and Two Corpus Obligations

59 A robot-learning episode must answer a practical question: what must be known before it can be
60 safely converted, replayed, synchronized, augmented, or used for evaluation? WorldEpisode answers
61 it with five coupled graphs at the level of a single episode record, two corpus-level obligations that
62 follow from those graphs, and one scale requirement for production datasets. Table 1 maps each
63 question to the structure that answers it, the requirement IDs that make it checkable, and the case in
64 Section 4 that tests it. WorldEpisode specializes these requirements for robot-learning episodes.

65 2.1 The five coupled graphs

66 Figure 1 shows the five graphs that make a spatial state replayable and auditable rather than just
67 stored. Making them mutually referential is what turns a silent error into a detectable one: an identity,
68 timing, role, event, or lineage error becomes a broken reference instead of a hidden assumption.




Figure 1: The five-graph model. WorldEpisode makes identity, frames/clocks, state roles, temporal events, and provenance mutually referential so conversion, synchronization, and replay errors become detectable instead of hidden.

69 **World revision and deltas** (WORLD.001, TRACE.001). Each record references one immutable,
70 content-addressed world revision plus an ordered list of deltas; any change to geometry, scale,
71 calibration, collision proxies, physical parameters, or object identity creates a new revision. Content
72 hashes let two datasets prove they refer to the same world without relying on filenames.

73 **Identity graph** (ENTITY.001). Every physical or logical entity receives a persistent identifier,
74 and the same `entity_id` must be used by segmentation masks, Gaussian groups, meshes, collision
75 bodies, simulator actors, language references, and interaction events. If a language command says
76 “pick up the mug,” the perception mask, the simulator object, and the collision proxy must identify
77 the same mug.

78 **Frame and clock graphs** (FRAME.001, TIME.001, ACTION.002). Every spatial value declares
79 its frames, units, transform direction, and validity interval, consistent with and extending ROS
80 conventions [16]; every stream declares a clock domain, timestamps, synchronization mapping, and
81 estimated drift, with the rigor NCore demonstrates for pose graphs [15]. Actions declare command
82 time and effective motor time rather than assuming those are the same instant. This graph turns a
83 suspicious replay from “the policy is bad” into a diagnosable offset, transform-direction, or unit error.

84 **Representation-role graph** (REP.001). One entity can have many representations (splat, mesh,
85 point cloud, collision proxy, mass/inertia record), each declaring a role and a digest-verified asset
86 descriptor. The role field is essential: a splat may be valid for appearance but invalid for collision,
87 while a simplified proxy may be valid for collision but unsuitable for visual training.

88 **Temporal-event graph** (TRACE.001). Records store interaction events — contact, grasp, attach-
89 ment, collision patch, reset, intervention, failure — each referencing entities, time, confidence, and
90 source, turning raw coordinate streams into auditable milestones.

91 **Provenance graph** (PROV.001). Every derived asset identifies its source episodes, sensors, recon-
92 struction algorithm and version, calibration revision, hashes, creator, license, and quality metrics,
93 mapping to existing provenance vocabularies rather than inventing an incompatible one.

94 The full normative field lists are specified in the v0.1 RFC and its JSON Schemas; the paper keeps
95 only what each graph must guarantee.

96 2.2 Two corpus-level obligations, plus scale

97 **Loss-explicit conversion** (CONVERT.001). WorldEpisode reuses existing standards rather than
98 replacing them. USD [2] carries composed simulation worlds, glTF carries Gaussian appearance
99 assets [19], NCore carries capture and pose-graph conventions [15], Rerun carries physical-data
100 recordings with column chunks [18], ROS/MCAP carries logs, and domain-native engines carry
101 their own telemetry. A converter may externalize or approximate fields, but it must emit a machine-
102 readable loss report instead of silently discarding semantics. The standard permits lossy conversion;
103 it forbids silent lossy conversion.

104 **Lineage-safe splits** (SPLIT.001). Record-level random splits are not enough when multiple records
105 share a reconstructed room, map patch, physical object, source capture, calibration, generated asset
106 lineage, or state ancestry — exactly what the provenance graph records. A WorldEpisode split mani-
107 fest can therefore enforce disjointness by constraints such as `world_lineage`, `physical_entity`,
108 `source_capture`, `reconstruction_run`, or `generated_asset_family`.

109 **Scale** (DATASET.001). Large robot corpora require a dataset manifest above the episode records:
110 namespaces, resolver priority, shard catalogs, materialized indexes, split catalogs, and append-only
111 versions. This lets a reader locate all SO-101 pick-place traces for a world-disjoint split by index
112 lookup rather than by crawling storage paths.

113 The initial WorldEpisode profile is intentionally narrow: rigid tabletop manipulation with fixed-
114 base single- or dual-arm robots. That bounded domain keeps conformance executable rather than
115 aspirational. Future profiles can extend this specification to humanoids, locomotion, deformables,
116 fluids, multi-agent game worlds, and autonomous-driving logs.

117 3 The Sidecar Contract, Made Executable

118 WorldEpisode delivers the contract as a sidecar blueprint around existing formats. It does not ask a
119 LeRobot dataset to become a USD scene, or a USD scene to become a policy-training table. Instead,
120 it records the relationships those containers need in order to agree about the same state: the state
121 revision, persistent entities, space/time conventions, transition contract, asset roles, interaction events,
122 provenance, and split lineage. The sidecar is not a mandate to package a global robot corpus as
123 millions of small folders: the single-episode record is the semantic unit, and production datasets add
124 a manifest layer (namespaces, URI resolvers, shard catalogs, materialized indexes, split manifests,
125 append-only snapshots) so that opening a corpus starts from the manifest and its indexes rather than
126 from a recursive storage scan.

127 Native containers continue to store the arrays, video, logs, worlds, and assets for which they were
128 designed. The sidecar stores only the relationships otherwise scattered across loader assumptions,
129 making cross-container identity, timing, role, lineage, and conversion loss checkable.

130 3.1 Runtime adapters

131 WorldEpisode is not tied to a single rendering or physics engine. A runtime becomes a trusted
132 target by implementing an adapter that preserves the same semantic invariants and emits replay
133 evidence, which keeps the integration direction clean: datasets are not rewritten around simulator-
134 native assumptions; simulators compile their scene loading, action interfaces, and replay reports
135 against the sidecar. The adapter contract has three layers:

- 136 • **Invariant interface**: ingests immutable world revisions, frame and clock graphs, entity identity,
137 representation roles, action channels, asset digests, provenance, and quality records.
- 138 • **Asynchronous schema extension**: lets an engine register cloth, deformable, fluid, articulation,
139 material, or procedural-scene extensions without weakening the core episode record.
- 140 • **Deterministic replay accountability**: records runtime identity, version, solver, timestep, actuator
141 parameters, initialization state, command/effective timing, latency, tolerance envelopes, and
142 measured drift.

The current implementation includes tested minimal MuJoCo [20] and Genesis [6] replay adapters for the same LeRobot control-loop trace. The Isaac mapping is emitted but untested, and SAPIEN remains adapter-required. The claim is therefore not simulator-independent physics but runtime-neutral accountability: every simulator reports against the same state contract, so cross-runtime drift becomes measurable rather than hidden inside loader conventions.

3.2 The executable robotics profile

Complete action semantics. An action vector is not portable without its operational contract. Each action channel must declare actuator, control mode, parameterization, reference frame, units, absolute/delta/velocity semantics, limits, command rate, command timestamp, effective timestamp, latency model, interpolation, normalization, gripper semantics, and missing-value policy. A vector such as $[0.01, 0, 0, 0, 0, 0, 1]$ is otherwise ambiguous: it could be a tool-frame delta, base-frame delta, velocity command, normalized policy output, or delayed target increment.

Implementation artifact. We implement the contract as four executable components: a JSON Schema for structural validity, a machine-readable conformance requirement file, an installable semantic validator, and a round-trip experiment runner. The validator maps each diagnostic to a stable requirement identifier such as `TIME.001`, `FRAME.002`, or `ASSET.002`. The experiment runner takes a `WorldEpisode` package, exports binding-native artifacts plus sidecars, reimports them, compares a versioned semantic projection, injects controlled faults, and writes reproducible result artifacts. The projection itself is published as `conformance/projections/uss-core-23.v0.json` and validated against `schemas/semantic-projection-v0.schema.json`. The contract is represented independently in the public schema and validator rather than only in its initial implementation.

One-line preflight. The validator is packaged as `worldepisode`. A blocking preflight can be run from the shell with `worldepisode preflight <dataset-or-manifest>` or from a training script with a single Python call such as `preflight_learner(dataset.root).raise_if_failed()`. Native LeRobot folders and `Rerun .rrd` recordings are recognized directly; without a `WorldEpisode` manifest or sidecar, they fail closed on missing world revision, entity identity, representation role, action timing, frame/clock, split-lineage, and conversion-loss controls. This makes the semantic audit a pre-training check rather than a separate post-hoc documentation exercise. The repository pairs the validator with converters, leakage and replay tools, a clean-room reader, and a runtime adapter report; Section 4 presents what each one measures.

4 Evidence: What the Audit Finds

We test what drift changes (Cases 1–3), what the contract detects or preserves (Cases 4–6), and which controls remain unavailable in 5 public benchmarks (Case 7). All measured data are robotics. Appendix A records claim boundaries and Appendix E records unexecuted gates.

4.1 What drift costs

Case 1: task-scene proxy leakage in evaluation splits (SPLIT.001). We audit 400 teleoperated episodes from ArmnetBench v0.1’s SO-101 LeRobot release [3]. Public metadata does not expose physical scene IDs, so we hash source revision, robot type, task ID/text, and camera schema into 8 *task-scene proxy* lineages. A seeded 80/20 episode split (320/80 episodes) puts every test proxy in training. The comparison split holds out task IDs 3 and 4 (300/100 episodes), yielding 0.000 proxy overlap. *This is also a task-disjoint shift*; it cannot identify the causal effect of scene overlap.

A 2-layer Torch MLP (64 units/layer, AdamW, 12 epochs, 5 seeds) maps state, episode phase, timestamp, and task one-hot features to action. The primary metric, mean episode nRMSE, changes from 0.183 (± 0.001 seed SD; 95% crossed seed-episode bootstrap CI [0.171, 0.196]) to 0.381 (± 0.020 ; [0.364, 0.399]); the increase is 0.197 [0.177, 0.219]. As a secondary diagnostic, the rate below a fixed 0.25 nRMSE threshold is 0.853 (± 0.016 seed SD) versus 0.000 (± 0.000). This threshold was not selected on a validation set and is not task success. Task one-hot features make this probe particularly sensitive to the task shift.

A deterministic ridge policy using only a 3-frame state history has mean episode nRMSE 0.157 [0.145, 0.170] versus 0.255 [0.247, 0.262]; the increase is 0.097 [0.082, 0.112]. Its secondary thresholded rates are 0.925 (74/80; [0.846, 0.965]) versus 0.420 (42/100; [0.328, 0.518]); its independently bootstrapped rate difference is 0.505 [0.390, 0.615]. We use 20,000 resamples. MLP intervals cross optimization seeds with held-out episodes, pair sampled seeds across protocols, and resample the different episode sets independently; ridge intervals resample episodes only. Thus the experiment establishes that the random split is not proxy-lineage-disjoint and that evaluation changes materially under a harder task-scene holdout; it does not establish scene-only leakage, policy rollout impact, or benchmark inflation.

The artifact includes digest-verified split packages and ACT/Diffusion job specifications. A pinned LeRobot 0.6.0 compatibility probe loads the representative package, but 0/2 policy paths complete a training step: all 2 stop during model construction because the low-dimensional package has neither an observation image nor an environment-state feature. We do not relabel joint proprioception to bypass that input contract. Policy results and rollouts therefore remain open (Appendix E).

Case 2: action-timing and replay drift (ACTION.002). LeRobot’s control stack can decouple policy inference from the fixed-rate low-level controller: commands may be queued, chunked, interpolated, or consumed after policy output [4]. We estimate an integer action/state lag from 320 calibration trajectories spanning 8 tasks, freeze it, and test 80 source-episode-disjoint trajectories with 0 source-episode overlap. The selected lag is 3 frames (150 ms). On identical within-episode target rows, zero-delay alignment has pooled RMSE 4.845 source position units and frozen-delay alignment has 1.934. The mean paired episode improvement is 3.087 (95% CI [2.890, 3.286]); all 80/80 held-out trajectories improve.

This is an action/state telemetry-lag proxy, not measured motor latency. The source exposes 1 float32 frame-timestamp column, no motor-effective timestamps, 1 robot type, and 1 controller configuration. Using the derived time shift, timestamp zero-order hold gives 2.021 RMSE whereas linear interpolation gives 1.934, showing that the timestamp comparison is interpolation-sensitive rather than independent queue evidence.

Separately, 1 trajectory has 303 6-joint samples at 30 Hz; its calibration prefix selects 4 frames, or 133 ms, and its held-out suffix changes alignment RMSE from 4.732 to 1.862 degrees. A WorldEpisode contract declares the estimated command/effective semantics, interpolation, and latency model. We execute it through 2 minimal position-servo adapters: MuJoCo changes from 3.425 to 1.563 degrees, while Genesis same-trace replay changes from 3.425 to 1.563. Their naive values (3.4247237 and 3.4247236) are nearly identical because both integrate the same minimal servo model, not because their physics agree. The Isaac adapter is emitted as a ready mapping from the same contract, but is not tested in this environment.

A dependency-free replay-adapter conformance harness guards additional runtime adapters: across 3 scheduling cases, including constant frame delay, zero-order hold for missing commands, and asynchronous queue selection, a naive scheduler fails while the contract-aware scheduler matches the expected trace with zero numerical error (0.0 maximum RMSE). This validates scheduler and same-trace adapter semantics only; it does not claim contact-rich task replay, Isaac/SAPIEN execution, or simulator-independent behavior.

Case 3: controlled real-to-sim contract drift (ACTION.001, REP.001). 2 deterministic proxies isolate an absolute-radian versus delta-degree controller mismatch and an appearance-only Gaussian export that drops collision geometry. Drifted contracts produce 2/2 proxy simulation successes and 0/2 deployment-proxy successes; explicit action and representation-role contracts recover 2/2. This is a mechanism check, not a hardware or RoboSnap rerun.

Figure 2 summarizes the measured changes for the displayed experiment families where explicit state semantics alter the outcome.

4.2 What the contract catches

Case 4: validator fault detection (all requirement IDs). We inject 14 faults into the baseline package: missing clock domains, incomplete cross-clock mappings, unknown frames, missing transform intervals, unknown event entities, missing representation roles, digest errors, action-




Figure 2: Summary of the displayed audits. The proxy-overlap result is confounded with task holdout; replay uses 1 trajectory; sidecar retention is by construction.

contract omissions, non-content-addressed worlds, missing delta sequences, missing provenance, and missing quality records. The validator detects all 14 expected requirement failures with 1.000 recall and 0.933 precision. 2 hand-authored independent fixtures also trigger the expected diagnostics. The single false-positive requirement is useful rather than harmful: removing an action reference frame triggers both ACTION.001 and the downstream FRAME.001 reference check. The same suite includes a unit sanity check from lerobot/svla_so101_pickplace: the first mirrored SO-101 frame has 6 joint values with maximum magnitude 98.924 in degrees and 1.727 after radian conversion, so treating those degree values as radians would violate the action/state unit contract even though the tensor shape is valid.

A pilot natural-source audit records 19 source omissions or evaluation risks across 2 active LeRobot artifacts plus source-level DROID and BridgeData V2 metadata. These are not maintainer-confirmed bugs or prevalence estimates. A 4-case preflight regression passes a valid sidecar and rejects an invalid manifest plus native LeRobot/Rerun artifacts that lack sidecars. An internal reader independent of the SDK path detects all 18 expected requirement instances across 17 fixtures; this tests artifact readability, not external adoption.

Case 5: loss-explicit dataset conversion (CONVERT.001). To test an actual LeRobotDataset rather than only a capability model, we implement LeRobotDataset v3 $\rightarrow$ WorldEpisode $\rightarrow$ LeRobotDataset v3 on 5-episode batches from 2 pinned public datasets: lerobot/svla_so101_pickplace [11] and lerobot/pusht [10]. The converter downloads the real Parquet metadata and data shards, builds a schema-valid WorldEpisode manifest and sidecar, exports a compact LeRobotDataset v3 package, reads it back, and asserts exact equality. These 10 committed episode packages provide inspectable per-episode evidence.

For scale, we convert every episode assigned to one immutable source Parquet shard from each of 3 datasets: the two above and ArmnetBench v0.1 [3]. No episode-level sampling is used. Across 271 episodes, including 2 multi-camera datasets, the round trip preserves 43,601 source-native paired action/state samples, indices, sample timestamps, and video timestamp ranges with maximum absolute error 0.0. The 5.2 MiB input produces 17.5 MiB of temporary packages in 48.24 s with 393.3 MiB peak resident memory. Source video payloads are neither downloaded nor converted; only their stream and timestamp metadata are audited. This is a complete-selected-shard measurement, not full-corpus throughput.

Both audits preserve derived physical-frame records through the sidecar and do not fake missing semantics. The same 4 source-absent fields—camera extrinsics, robot/world calibration, action units, and controller latency—are reported as conversion loss rather than guessed.

Case 6: binding retention (CONVERT.001). 7 pilot bindings are compared against a versioned 23-field projection. Non-reference native retention ranges from 0.174 to 0.391; sidecars recover 1.000 for tested dataset/log/world bindings by serializing the missing fields. This is a round-trip check and gap inventory, not a universal format ranking; Appendix Table 3 reports the per-binding values.

4.3 Scaling the audit

Case 7: source-level benchmark audit (SPLIT.001). We audit public metadata for Open X-Embodiment/RT-X [17], DROID [9], BridgeData V2 [21], LIBERO [12], and CALVIN [13]. The audit asks whether public metadata exposes lineage splits, content-addressed world revisions, persistent entity identity, action timing contracts, frame/clock graphs, and loss reports. Each has at least one high-severity open control. A DROID-100 proxy rerun over 100 episodes and 32,212 frames scores 0.516 under its corrected proxy split versus 0.496 randomly, so it provides no inflation evidence. The proof gate therefore records 0 famous-benchmark inflation claims.

The same audit protocol also runs above the episode level: a dataset-scale manifest audit and a generated catalog benchmark (Appendix F) make the DATASET.001 scale claim executable on catalog-side evidence only.

5 Positioning

A contribution of falsifiable invariants. The experiments show several ways a result can be wrong while the files are valid: a split can leak world lineage, replay can drift because control timing is underspecified, conversion can discard physical semantics that the target container has no native place to store, and a live virtual or AV state can drift from its authoritative revision. WorldEpisode is useful only because these failures are observable and testable. The sidecar is the mechanism; the scientific claim is that explicit spatial-state semantics change what an embodied-agent result means. A container adapter either preserves those invariants, externalizes them with declared loss, or fails conformance.

A small semantic core with multiple bindings. The bindings are evaluated at very different depths: LeRobotDataset v3 is the only binding exercised with active real-data experiments (conversion, leakage, replay timing); Rerun, NCore, MCAP/ROS, OpenUSD/SimReady, and glTF/SPZ are covered only as binding-retention pilots (Table 3). OpenUSD standardizes how the 3D world is composed; WorldEpisode standardizes how a robot modifies state within that space over time without silent data corruption. The same invariant structure should generalize to a game-engine or AV profile, though only the robotics profile is implemented and evaluated here.

Real-to-sim pipelines and benchmarks are the users, not the competitors. Gaussian-splat and OpenUSD real-to-sim pipelines are the systems that most need a state and episode contract: a reconstructed room can be excellent visual training data and still be unsafe for policy evaluation if collision roles are missing, lineage leaks across splits, or the replayed action vector is interpreted under the wrong hardware contract. The same stance governs benchmark auditing: when public metadata cannot answer whether train and test share a scene lineage, asset family, source capture, or timing convention, WorldEpisode marks the published score as unaudited with respect to that control — and marks it inflated only after the same policy protocol is rerun under the corrected split or timing contract.

Standards alignment. The project aligns with standards bodies rather than antagonizing them. Static 3D robot-map exchange, humanoid dataset lifecycle, OpenLABEL annotation, glTF Gaussian splats, USD composition, and NCore capture conventions are all useful neighbors. WorldEpisode is the executable profile and bridge across those layers.

6 Limitations and Open Gates

The current evidence is deliberately scoped: it exposes hidden failure mechanisms rather than claiming full robot-learning, game-engine, or autonomous-driving generality. The boundary of every empirical claim is tabulated in Appendix A, and each open gate’s commands and acceptance rules are boxed in Appendix E; this section states what matters most.

Beyond robotics (future work). All measured evidence in this paper is the WorldEpisode robotics profile; the broader Universal Spatial State framing is proposed, not demonstrated outside robotics. 2 deterministic pilots only illustrate that the same vocabulary can express non-robotics drift. In a

game-engine collision pilot, a client loads a stale arena collision asset after an authoritative patch, which the asset-digest, state-revision, and role checks catch. In an autonomous-driving clock pilot, camera and lidar logs deserialize correctly, but a 50 ms undeclared clock offset at 15 m/s creates a 0.75 m fusion error against a 0.20 m tolerance, which the declared clock mapping corrects to 0.0 m. These are illustrative pilots, not production game-engine or AV benchmark results. A stronger paper should add at least one public non-robotics corpus and an independently maintained adapter.

Principal limitations. The headline holdout is confounded: task identity is part of the proxy-lineage key, so it measures a task–scene shift and cannot attribute the change to scene overlap alone. The MLP uses 5 optimization seeds and task one-hot features; crossed bootstrap intervals capture seed and held-out-episode variation, but 5 seeds remain a small training sample and the 0.25 nRMSE threshold is reported only as a secondary diagnostic rather than task success. ACT/Diffusion training and rollout evidence remain open; the pinned compatibility probe stops both policy paths before training because source images were not materialized. The timing audit calibrates on 320 trajectories and evaluates 80 held-out trajectories, but it has no motor-effective timestamps and covers 1 robot type and 1 controller configuration. MuJoCo and Genesis still execute the same minimal servo model on 1 trajectory, not independent contact-rich physics. Conversion covers complete selected Parquet shards from 3 datasets, not their full corpora or source video payloads. Injected faults establish detection on known mutations, not natural prevalence or false-positive rates in deployment.

The implementation is also single-sourced. The clean-room reader is internal, the package is not yet an upstream LeRobot/Rerun integration, and no external compatible dataset is available. The 5-dataset natural-source audit mixes 2 active artifacts with 2 source-level metadata reviews and has no maintainer adjudication.

Societal impact. Fail-closed state checks may reduce unsafe or irreproducible robot evaluations and wasted training compute. Conversely, an incomplete contract can create false assurance, provenance fields can expose sensitive capture locations or operators, and content-addressed lineage can aid dataset linkage. Deployments should minimize or access-control provenance, retain native safety interlocks, and treat conformance as an audit aid rather than a safety certificate.

7 Conclusion

Robot-learning failures often arrive disguised as clean data: valid tensors, valid videos, valid scene files, and invalid state assumptions. WorldEpisode makes those assumptions explicit. It binds interaction data to immutable state revisions, gives persistent identity to entities, declares space/time/transition semantics, assigns roles to assets, records events and provenance, and reports conversion loss instead of hiding it.

In the reference experiments, a random split overlaps every task–scene proxy lineage; a task-confounded holdout increases mean episode nRMSE from 0.183 to 0.381 for the MLP and from 0.157 to 0.255 for deterministic temporal ridge. This diagnoses an invalid random split but does not isolate scene leakage. Modeling a 4-frame action delay reduces held-out alignment error on 1 SO-101 trajectory, and the validator catches every injected fault class. The contribution is the auditable contract and its failure boundaries, not a claim of physical rollout improvement.

The next experiments must disentangle task from physical scene lineage, run ACT/Diffusion and rollouts over matched seeds, test delay on multiple trajectories, and add contact-rich replay. Until then, the paper makes no famous-benchmark inflation, simulator-independent physics, or mature-standard claim.

References

- [1] Mubashara Akhtar, Omar Benjelloun, Costanza Conforti, Luca Foschini, Pieter Gijssbers, Joan Giner-Miguel, Sujata Goswami, Nitisha Jain, Michalis Karamousadakis, Satyapriya Krishna, Michael Kuchnik, Sylvain Lesage, Quentin Lhoest, Pierre Marcenac, Manil Maskey, Peter Mattson, Luis Oala, Hamidah Oderinwale, Pierre Ruysen, Tim Santos, Rajat Shinde, Elena Simperl, Arjun Suresh, Goeffry Thomas, Slava Tykhonov, Joaquin Vanschoren, Susheel Varma, Jos van der Velde, Steffen Vogler, Carole-Jean Wu, and Luyao Zhang. Croissant: A metadata format for ML-ready datasets. In *Advances in*

- 380 *Neural Information Processing Systems*, volume 37, pages 82133–82148, 2024. doi:
381 10.52202/079017-2610. URL [https://proceedings.neurips.cc/paper_files/
382 paper/2024/file/9547b09b722f2948ff3ddb5d86002bc0-Paper-Datasets_and_
383 Benchmarks_Track.pdf](https://proceedings.neurips.cc/paper_files/paper/2024/file/9547b09b722f2948ff3ddb5d86002bc0-Paper-Datasets_and_Benchmarks_Track.pdf).
- 384 [2] Alliance for OpenUSD. OpenUSD: Universal scene description. <https://openusd.org/>,
385 2026. Accessed 2026-07-13.
- 386 [3] Armnet. ArmnetBench v0.1: LeRobot single-arm SO-101 release. [https://
387 huggingface.co/datasets/armnet/armnetbench_v01_lerobot_so101](https://huggingface.co/datasets/armnet/armnetbench_v01_lerobot_so101), 2026. Pinned
388 revision 2e5e89aee0e7f081078d9d6ab3b279fc4b83ea84; accessed 2026-07-13.
- 389 [4] Remi Cadene, Simon Aliberts, Francesco Capuano, Michel Aractingi, Adil Zouitine, Pepijn
390 Kooijmans, Jade Choghari, Martino Russi, Caroline Pascal, Steven Palma, Mustafa Shukor,
391 Jess Moss, Alexander Soare, Dana Aubakirova, Quentin Lhoest, Quentin Gallouedec, and
392 Thomas Wolf. LeRobot: An open-source library for end-to-end robot learning. In *International
393 Conference on Learning Representations*, 2026. arXiv:2602.22818.
- 394 [5] Timnit Gebru, Jamie Morgenstern, Briana Vecchione, Jennifer Wortman Vaughan, Hanna
395 Wallach, Hal Daumé III, and Kate Crawford. Datasheets for datasets. *Communications of
396 the ACM*, 64(12):86–92, 2021. doi: 10.1145/3458723. URL [https://doi.org/10.1145/
397 3458723](https://doi.org/10.1145/3458723).
- 398 [6] Genesis Authors. Genesis: A generative and universal physics engine for robotics and beyond.
399 <https://github.com/Genesis-Embodied-AI/Genesis>, 2024. Version 0.3.3 used in the
400 replay experiment; accessed 2026-07-20.
- 401 [7] Hugging Face. LeRobotDataset v3.0. [https://huggingface.co/docs/lerobot/
402 lerobot-dataset-v3](https://huggingface.co/docs/lerobot/lerobot-dataset-v3), 2026. Accessed 2026-07-13.
- 403 [8] Sayash Kapoor and Arvind Narayanan. Leakage and the reproducibility crisis in machine-
404 learning-based science. *Patterns*, 4(9):100804, 2023. doi: 10.1016/j.patter.2023.100804. URL
405 <https://doi.org/10.1016/j.patter.2023.100804>.
- 406 [9] Alexander Khazatsky, Karl Pertsch, Suraj Nair, Ashwin Balakrishna, Sudeep Dasari, Siddharth
407 Karamcheti, Soroush Nasiriany, Mohan Kumar Srirama, et al. DROID: A large-scale in-the-
408 wild robot manipulation dataset. In *Robotics: Science and Systems*, 2024. URL [https:
409 //www.roboticsproceedings.org/rss20/p120.html](https://www.roboticsproceedings.org/rss20/p120.html). Dataset released under CC-BY-
410 4.0.
- 411 [10] LeRobot. pusht: A public LeRobotDataset v3 dataset. [https://
412 huggingface.co/datasets/lerobot/pusht](https://huggingface.co/datasets/lerobot/pusht), 2025. Pinned revision
413 7628202a2180972f291ba1bc6723834921e72c19; accessed 2026-07-13.
- 414 [11] LeRobot. svla_so101_pickplace: A public LeRobotDataset v3 dataset. [https://
415 huggingface.co/datasets/lerobot/svla_so101_pickplace](https://huggingface.co/datasets/lerobot/svla_so101_pickplace), 2025. Pinned revision
416 f641879e22172be7e8161d5e6c1503c2d2feb657; accessed 2026-07-13.
- 417 [12] Bo Liu, Yifeng Zhu, Chongkai Gao, Yihao Feng, Qiang Liu, Yuke Zhu, and Pe-
418 ter Stone. LIBERO: Benchmarking knowledge transfer for lifelong robot learn-
419 ing. In *Advances in Neural Information Processing Systems*, volume 36, 2023.
420 URL [https://proceedings.neurips.cc/paper_files/paper/2023/hash/
421 8c3c666820ea055a77726d66fc7d447f-Abstract-Datasets_and_Benchmarks.html](https://proceedings.neurips.cc/paper_files/paper/2023/hash/8c3c666820ea055a77726d66fc7d447f-Abstract-Datasets_and_Benchmarks.html).
- 422 [13] Oier Mees, Lukas Hermann, Erick Rosete-Beas, and Wolfram Burgard. CALVIN: A benchmark
423 for language-conditioned policy learning for long-horizon robot manipulation tasks. *IEEE
424 Robotics and Automation Letters*, 7(3):7327–7334, 2022. doi: 10.1109/LRA.2022.3180108.
425 URL <https://doi.org/10.1109/LRA.2022.3180108>.
- 426 [14] Margaret Mitchell, Simone Wu, Andrew Zaldivar, Parker Barnes, Lucy Vasserman, Ben Hutchin-
427 son, Elena Spitzer, Inioluwa Deborah Raji, and Timnit Gebru. Model cards for model reporting.
428 In *Proceedings of the Conference on Fairness, Accountability, and Transparency*, pages 220–
429 229, 2019. doi: 10.1145/3287560.3287596. URL [https://doi.org/10.1145/3287560.
430 3287596](https://doi.org/10.1145/3287560.3287596).

- 431 [15] NVIDIA. NCore Specification: Data Conventions. [https://nvidia.github.io/ncore/](https://nvidia.github.io/ncore/data/conventions)
432 [data/conventions](https://nvidia.github.io/ncore/data/conventions), 2026. Accessed 2026-07-13.
- 433 [16] Open Robotics. ROS REP 103: Standard units of measure and coordinate conventions. [https://](https://www.ros.org/reps/rep-0103.html)
434 www.ros.org/reps/rep-0103.html, 2026. Accessed 2026-07-13.
- 435 [17] Open X-Embodiment Collaboration. Open X-Embodiment: Robotic learning datasets and rt-x
436 models. In *IEEE International Conference on Robotics and Automation*, pages 6892–6903, 2024.
437 doi: 10.1109/ICRA57147.2024.10611477. URL [https://doi.org/10.1109/ICRA57147.](https://doi.org/10.1109/ICRA57147.2024.10611477)
438 [2024.10611477](https://doi.org/10.1109/ICRA57147.2024.10611477).
- 439 [18] Rerun. A new data layer for robot learning. [https://rerun.io/blog/](https://rerun.io/blog/data-layer-for-robot-learning)
440 [data-layer-for-robot-learning](https://rerun.io/blog/data-layer-for-robot-learning), 2026. Accessed 2026-07-13.
- 441 [19] The Khronos Group. KHR_gaussian_splatting. [https://github.com/KhronosGroup/](https://github.com/KhronosGroup/glTF/blob/main/extensions/2.0/Khronos/KHR_gaussian_splatting/README.md)
442 [glTF/blob/main/extensions/2.0/Khronos/KHR_gaussian_splatting/README.md](https://github.com/KhronosGroup/glTF/blob/main/extensions/2.0/Khronos/KHR_gaussian_splatting/README.md),
443 2026. Accessed 2026-07-13.
- 444 [20] Emanuel Todorov, Tom Erez, and Yuval Tassa. MuJoCo: A physics engine for model-based
445 control. In *IEEE/RSJ International Conference on Intelligent Robots and Systems*, pages 5026–
446 5033, 2012. doi: 10.1109/IROS.2012.6386109. URL [https://doi.org/10.1109/IROS.](https://doi.org/10.1109/IROS.2012.6386109)
447 [2012.6386109](https://doi.org/10.1109/IROS.2012.6386109).
- 448 [21] Homer Rich Walke, Kevin Black, Tony Z. Zhao, Quan Vuong, Chongyi Zheng, Philippe Hansen-
449 Estruch, Andre Wang He, Vivek Myers, Moo Jin Kim, Max Du, Abraham Lee, Kuan Fang,
450 Chelsea Finn, and Sergey Levine. BridgeData V2: A dataset for robot learning at scale. In
451 *Proceedings of the 7th Conference on Robot Learning*, volume 229 of *Proceedings of Machine*
452 *Learning Research*, pages 1723–1736, 2023. URL [https://proceedings.mlr.press/](https://proceedings.mlr.press/v229/walke23a.html)
453 [v229/walke23a.html](https://proceedings.mlr.press/v229/walke23a.html). Dataset released under CC-BY-4.0.

A Evidence Boundaries

Table 2 states the boundary of each empirical claim in Section 4.

Table 2: What the current evidence supports, and what it does not yet support.

Claim area	Current evidence	Boundary
Leakage	Public ArmnetBench LeRobot audit with 400 teleoperated reference episodes, an executable Torch BC probe, a measured temporal ridge state/action baseline, an ACT/Diffusion gate harness, and 4 compact digest-verified LeRobot state/action split packages totaling 241,470 frames	Lineage is a task–scene proxy and the holdout is task-disjoint; scene-only leakage is not identified. The LeRobot 0.6.0 preflight loads the representative package, but 0/2 policy paths complete a training step because split packages omit the required image/environment-state modality; rollouts are not run
Conversion	Complete pinned source Parquet shards from 3 public LeRobotDataset v3 datasets: 271 episodes and 43,601 paired state/action samples with exact tensor, index, and timestamp equality	Selected Parquet shards rather than full corpora; source video payloads are not converted or used for a throughput claim
Replay timing	Lag calibration on 320 SO-101 trajectories, 80 source-episode-disjoint evaluations across 8 tasks, tested same-trace MuJoCo and Genesis position-servo adapters, and scheduler conformance	Telemetry proxy with no motor-effective timestamps, 1 robot type, and 1 controller configuration; simulator adapters use 1 trajectory, with no contact-rich task rollout, Isaac result, or SAPIEN result
Meta-simulator neutrality	Adapter contract matrix over MuJoCo, Isaac Sim, Genesis, and SAPIEN plus same-trace MuJoCo/Genesis replay evidence	MuJoCo and Genesis are tested for the minimal replay profile; Isaac and SAPIEN are not replay-tested in this paper, and equal physics is not claimed
Generalization beyond robotics	Deterministic game-engine collision-patch and autonomous-driving clock-domain pilots using the same invariant vocabulary	Not measured on production game engines, public AV datasets, or fleet logs
Validation	14 injected requirement faults, 2 independent hand-authored fixtures, and a pilot natural-source corpus over 5 public datasets	The 5-dataset count is met through active LeRobot artifacts plus source-level public benchmark metadata; dataset-specific diagnostic reports now cover all cases. Maintainer validation: <i>Not defined yet</i>
Binding retention	Versioned <code>uss-core-23</code> semantic projection checked by executable artifacts	Pilot projection; not a universal score of each storage format
Famous benchmark call-out	Source-level audit over Open X-Embodiment, DROID, BridgeData V2, LIBERO, and CALVIN, plus a targeted DROID subset rerun tool and a separate inflation-proof gate	Missing public controls are not measured score inflation; the proof gate currently has 100-episode DROID subset evidence but 0 measured inflation claims
Real-to-sim drift	Controlled action-contract and representation-role ablations where drifted contracts succeed in simulation and fail under deployment proxies	Deterministic proxy; not a RoboSnap/DROID-Sim rerun or physical hardware deployment
Adoption	Public schema, installable local preflight package, validator, fixtures, governance files, and an internal clean-room reader that does not import the reference SDK	No PyPI release, upstream LeRobot/Rerun integration, external independent implementation, or external dataset release yet
Scale architecture	Dataset manifest schema, scalable example, executable audit, and a generated 32,768-shard catalog benchmark describing 1,073,741,824 episodes for open/index latency, partition pruning, digest-cache behavior, and resolver routing	Catalog-side benchmark only; no billion episode rows, payload bytes, distributed storage, concurrent readers, or institutional federation are measured

B Experiment and Compute Provenance

The supplementary artifact includes a machine-readable experiment manifest at `docs/experiments/experiment_manifest/experiment_manifest.json`. It joins 7 principal runs to pinned dataset revisions, split and code digests, effective configuration, seed policy, output digests, and the remote execution log. The measured environment provides an NVIDIA GB10 accelerator, 20 logical CPUs,

and 121.7 GiB RAM. The MLP run takes 46.69 s with 1510.7 MiB peak resident memory; the DROID and controlled-suite processes take 0.62 s and 3.84 s. The conversion-scale process takes 48.24 s with 393.3 MiB peak resident memory; its 2 preliminary failed runs are retained with their nonzero exit status and logs. The multi-trajectory timing audit takes 1.20 s with 247.6 MiB peak resident memory. The logs also record storage and exit status. Temporal ridge shares the controlled-suite process and is labeled with that accounting scope rather than assigned invented per-subtest compute.

The anonymous supplementary ZIP is built from an explicit allowlist with deterministic entry digests. Source-control history and local environments are excluded; text payloads are anonymized before hashing, binary payloads are scanned, and a separate audit checks extracted PDF text and metadata plus every ZIP path and payload.

C Third-Party Assets

The generated inventory at docs/experiments/third_party_asset_audit/asset_audit.json joins the experiment manifest to immutable dataset-card evidence. The 4 active datasets are licensed as follows: ArmnetBench under Apache-2.0, SVLA SO-101 pick-place under Apache-2.0, PushT under MIT, and the LeRobot DROID-100 mirror under MIT and CC-BY-4.0 (retaining both the mirror and upstream DROID obligations).

The supplement contains 48 source-derived Parquet files in round-trip and split packages, covered by 16 adjacent SOURCE_LICENSE.json records. It contains 0 source image or video files. Source-derived rows retain their source expression; WorldEpisode-authored manifests, diagnostics, and aggregate reports are CC0-1.0 unless stated otherwise. The generated THIRD_PARTY_ASSETS.md records attributions, immutable card digests, installed software licenses, and context-only datasets. Apache-2.0 and PushT MIT texts are distributed with the artifact.

The official NeurIPS 2026 style is pinned by digest. Because its distributed file contains no separate SPDX or license statement, the inventory records its license as NOASSERTION and states the author-kit submission use instead of inventing a license. This inventory is provenance documentation, not legal advice.

D Pilot Binding Retention

Table 3: Semantic field retention in the versioned 23-field pilot projection. The reference binding retains 1.000 natively and is omitted. Sidecar recovery is by construction.

Binding	Native	With WorldEpisode
LeRobot v3	0.261	1.000
Rerun RRD	0.348	1.000
NCore	0.304	1.000
MCAP/ROS2	0.261	1.000
OpenUSD/SimReady	0.391	1.000
glTF Gaussian asset	0.174	0.304

E Open Gates: Commands and Acceptance Rules

Open result, not claimed: results not claimed in this release

The artifact keeps unfinished claims executable instead of hiding them. The current RFC release gate is: `python3 tools/release_readiness.py -strict-rfc`. The full-standard gate is: `python3 tools/release_readiness.py -strict-full`. The latter is expected to fail until the two gates boxed below are backed by committed reports and, additionally, external standard evidence exists: maintainer-reviewed diagnostics, an independently written reader/exporter, and at least one externally published WorldEpisode-compatible dataset. The same gates are indexed in docs/experiments/open_reproduction_gates/open_reproduction_gates.json.

Open result, not claimed: ACT/Diffusion and rollout impact

Result: *Not defined yet*. Rollout impact: *Not defined yet*. Scene-only leakage: *Not defined yet*. First replace the task-derived proxy with physical scene/source-capture IDs and construct task-stratified random and scene-lineage-disjoint splits. Then regenerate split packages and LeRobot job specs and run ACT/Diffusion inside a LeRobot training environment. Before training, materialize source images or a semantically valid environment-state feature and require the compatibility preflight to complete a training step.

Required artifacts: a proof that task support is identical across protocols and scene lineage is disjoint, policy checkpoint digests, train/eval configs with at least five seeds, offline action metrics, rollout traces or videos with content digests, and an updated policy-gate report.

Acceptance rule: both split protocols must preserve task support, at least one ACT or Diffusion Policy must report both protocols over multiple seeds, and at least one simulator or physical rollout must use the same manifests. Running the current task-confounded jobs does not close this gate.

Open result, not claimed: famous-benchmark score-inflation proof

Result: *Not defined yet*. A targeted DROID subset rerun can be attempted with `uv run -with pyarrow -with requests -with numpy python tools/famous_benchmark_policy_rerun.py -benchmark droid_100 -required`. The proof contract is then enforced with `python3 tools/benchmark_inflation_gate.py -required`.

Required artifacts: benchmark-specific WorldEpisode conversion, lineage/timing audit, faithful published-protocol policy rerun, paired corrected evaluation under the same metric, and a measured baseline-minus-corrected score drop.

Acceptance rule: the gate must contain at least one inflation-proof valid rerun report with `measured_inflation=true`. Source-level metadata gaps alone are evidence of missing controls, not evidence that published scores are inflated.

F Dataset-Scale Audit (DATASET.001)

To address production scale, we audit the dataset manifest in `examples/scalable-corpus.worldepisode-dataset.json`. The audit validates 5 namespaces, 3 resolver schemes, 3 registries, 3 shards, 2 indexes, 1 append-only version, and 9 digest-addressed asset descriptors spanning Hugging Face, S3, and OCI URIs with local mirrors. It also checks that the manifest includes a world-lineage index, an asset-digest index, and a split-manifest shard. We then run a generated catalog benchmark with 32,768 trace shards describing 1,073,741,824 episodes. The committed report records the measured time for parsing and indexing the generated JSON catalog, 8 partition-pruning queries, digest-cache resolution, and resolver routing over all asset descriptors. This is catalog-side evidence only, but it makes the scale claim executable: opening a corpus starts from the manifest and indexes rather than from recursive storage discovery.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Justification: The abstract, Section 1, and Section 6 state the task–scene confound, one-trace replay scope, and open gates.

Guidelines:

- The answer [\[N/A\]](#) means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [\[No\]](#) or [\[N/A\]](#) answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.

- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Section 6 discusses the split confound, model and sampling uncertainty, replay and conversion scale, adoption, and societal risks.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper makes no theoretical claims and contains no theorems.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Section 4 reports pinned revisions, splits, sample counts, models, hyperparameters, metrics, and interval construction; the submitted artifact includes executable scripts and exact commands.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: The anonymized supplementary archive includes the validator, experiment scripts, compact split packages, committed reports, exact digests, and reproduction guide. An automated audit checks extracted PDF text and metadata plus all ZIP paths and payloads for identifying content.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not

including code, unless this is central to the contribution (e.g., for a new open-source benchmark).

- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Section 4 gives the split construction, feature sets, optimizer, epochs, seeds, sample counts, calibration split, primary continuous metrics, and secondary diagnostic thresholds.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [No]

Justification: Case 1 reports crossed seed–episode bootstrap intervals over 5 MLP seeds and episode intervals for deterministic ridge, but the 1-trajectory replay and exact conversion checks do not have repeated-run uncertainty estimates.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).

- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [No]

Justification: Appendix B and the supplementary experiment manifest record host, NVIDIA GB10 accelerator, CPU count, RAM, storage, wall time, peak memory, and exit status for the principal remote runs. The answer remains No because failed preliminary work and total-project compute are not comprehensively accounted.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [N/A]

Justification: The authors must review and attest to the NeurIPS 2027 Code of Ethics before submission; the 2027 code is not yet available.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: The societal-impact paragraph in Section 6 discusses safer evaluation and reduced compute waste alongside false assurance, privacy, and linkage risks.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to

generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.

- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The work releases no high-risk pretrained model or scraped dataset; it provides schemas, validators, fixtures, and derived audit reports.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: Appendix C names the licenses and pinned revisions for every active dataset. Each redistributed source-derived package carries machine-readable license evidence and attribution; no source media are redistributed. The supplementary artifact includes the generated inventory, license notices, software sources, and the official author-kit style's pinned digest and distribution context; no license is invented where upstream provides no separate assertion.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.

- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset’s creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: The submitted artifact documents the schema, validator, conformance fixtures, experiment scripts, governance, limitations, and Apache-2.0 reference-code license.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: This work performs no new crowdsourcing or human-subject study; it audits existing public robot-learning datasets.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: This work performs no new crowdsourcing or human-subject study.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.

831 • For initial submissions, do not include any information that would break anonymity (if
832 applicable), such as the institution conducting the review.

833 **16. Declaration of LLM usage**

834 Question: Does the paper describe the usage of LLMs if it is an important, original, or
835 non-standard component of the core methods in this research? Note that if the LLM is used
836 only for writing, editing, or formatting purposes and does *not* impact the core methodology,
837 scientific rigor, or originality of the research, declaration is not required.

838 Answer: [N/A]

839 Justification: No LLM or agent is an important, original, or non-standard component of the
840 contract, validator, experiment design, or reported measurements.

841 Guidelines:

842 • The answer [N/A] means that the core method development in this research does not
843 involve LLMs as any important, original, or non-standard components.

844 • Please refer to our LLM policy in the NeurIPS handbook for what should or should not
845 be described.