

Universal Spatial State: Diagnosing and Correcting Silent State Drift across Embodied and Virtual Agent Datasets

Alba Tellez Fernandez
URDF Studio

USS / WorldEpisode Contributors

Abstract

Valid spatial-data files can still describe the wrong experiment. A robot command stream may load successfully while being replayed at the wrong effective time; a train/test split may report strong policy performance while reusing the same physical scene lineage; a simulator asset may render correctly while referring to a stale state revision, frame convention, or digest. We call this failure mode *silent state drift*: structural validation passes, but the behavioral state being trained, replayed, or deployed no longer matches the state the experiment claims to evaluate.

We introduce *Universal Spatial State* (USS), a storage-, representation-, and runtime-neutral contract for auditing this drift. USS does not replace OpenUSD, LeRobotDataset, Rerun, MCAP, NCore, game telemetry, or autonomous-driving logs. Instead, it specifies the invariants those containers must preserve: persistent identity, immutable state revisions, explicit frame and clock graphs, representation roles, action and transition semantics, temporal events, provenance, uncertainty, asset digests, and lineage-aware splits. *WorldEpisode* is the robotics reference profile used to make USS executable on real LeRobot traces, reconstructed-world assets, conversion reports, and simulator adapters.

We evaluate WorldEpisode by converting common hidden assumptions into measurable failure tests. On ArmnetBench’s public SO-101 LeRobot release, random episode splitting leaks every tested world lineage and yields 0.850 offline behavioral-cloning success; a scene-disjoint split reduces leakage to zero (0.000) and the same probe to 0.000. On a real SO-101 trajectory, the inferred four-frame, 133 ms action delay reduces validation joint RMSE from 4.732 to 1.862 degrees when modeled explicitly. The same timestamp-aware contract reduces tested MuJoCo replay RMSE from 3.425 to 1.563 degrees and Genesis same-trace replay RMSE from 3.425 to 1.563 degrees. An active LeRobotDataset-to-WorldEpisode-to-LeRobotDataset converter preserves ten episodes across two public datasets, 1,935 state/action rows, timestamps, video timestamp ranges, and physical-frame records with maximum numerical error 0.0 while recording source-absent semantics as explicit loss. The validator detects all 14 injected fault classes and the independent fixture failures.

The contribution is a tested state-integrity layer for behavioral validity, not another monolithic file format. The current evidence remains deliberately scoped: ACT/Diffusion rollouts, famous benchmark score-inflation proofs, contact-rich multi-simulator task replay, Isaac/SAPIEN execution, external adoption, and production game or autonomous-driving audits are kept as open gates rather than claimed outcomes.

1 Introduction

Imagine training a behavioral-cloning policy for three weeks on a public robot dataset. The loss curves are smooth, the simulator reports an 85% success rate, and every Parquet, video, and scene file passes its local loader. Then the policy is deployed on the real arm and immediately misses the object. The failure is not a new neural architecture problem. The dataset silently encoded joint values in a unit the controller did not declare, the replay script treated command timestamps as

motor-effective timestamps, and the random split let the policy see the same reconstructed room in both training and evaluation. Nothing was corrupt, but the experiment was scientifically broken.

The same pattern appears outside robotics. A game client can load an old collision asset after a server patch, but the authoritative state now places an obstacle in the avatar path. An autonomous-driving replay can deserialize camera and lidar logs correctly while fusing samples from different clock domains. A cloud digital twin can point to the right-looking asset URI while using the wrong content digest. In each case the file is valid and the behavior is not. We call this failure class *silent state drift*.

This paper introduces Universal Spatial State (USS), a sidecar contract for making silent state drift auditable. USS is not a file format. It is a set of invariants that bind interactive spatial data to immutable state revisions: persistent identity, explicit space/time mappings, role semantics, transition contracts, temporal events, provenance, quality records, and lineage-aware splits. WorldEpisode is the robotics-heavy USS profile implemented and evaluated in this repository. Robotics is the deepest current stress test because it combines real hardware, asynchronous control, multi-camera data, reconstructed worlds, simulator exports, and policy evaluation.

Existing formats solve important pieces of this chain. The LeRobot library vertically integrates robot learning from motor interfaces through dataset collection, storage, streaming, policies, and deployment [3]. The current LeRobotDataset v3 layout stores high-frequency states, actions, and timestamps in Parquet, visual streams in MP4 shards, and episode indexing in metadata rather than filenames [6]. Rerun provides a physical-data layer and storage target for multimodal recordings. NCore documents explicit pose graphs and sensor conventions. OpenUSD, SimReady, glTF, and simulator-native formats describe composed worlds and assets. Gaussian splats are also moving into standard asset containers through glTF’s `KHR_gaussian_splatting` extension [14, 15]. Recent real-to-sim systems such as RoboSnap and DROID-Sim show that Gaussian-splat environments can support trajectory replay, synthetic data generation, and policy evaluation at dataset scale [13]. That progress makes the missing contract more urgent rather than less. A layered real-to-sim scene can be visually faithful while still dropping the collision role of an object, and a replayable trajectory can still be wrong if the policy action vector is interpreted under a different controller contract at deployment time.

The missing piece is not another monolithic file. A LeRobot table can store actions without knowing which immutable world revision they came from; a USD scene can store assets without knowing the policy action semantics that generated a trajectory; a ROS or MCAP log can store messages without knowing whether evaluation splits leak the same physical room. More generally, a game telemetry log or AV sensor recording can be locally valid without declaring the state revision, asset digest, clock mapping, or transition invariant that makes its behavior meaningful. USS fills that gap as a sidecar contract; WorldEpisode instantiates it for robot-learning episodes. OpenUSD standardizes how the 3D world is composed, but USS standardizes how any agent, whether a physical robot, a video game character, or an autonomous vehicle, modifies state within that space over time without silent data corruption.

USS is:

- **storage-neutral**: serializable through LeRobotDataset v3, Rerun, NCore, MCAP, game telemetry, AV logs, or a reference package;
- **representation-neutral**: compatible with splats, meshes, NeRFs, point clouds, collision proxies, sensor streams, and future assets;
- **runtime-neutral**: mappable to MuJoCo, Isaac Sim, SAPIEN, Genesis, game engines, AV replay systems, and future runtimes;

- **loss-explicit:** conversions may be lossy, but never silently lossy.

The paper’s core claim is therefore not a new monolithic file format. It is an executable spatial state contract: identity, frames, clocks, transition contracts, representation roles, events, provenance, uncertainty, and lineage-safe evaluation splits. The current artifact includes a JSON Schema, semantic validator, round-trip binding artifacts, an active public LeRobotDataset v3 conversion experiment, hand-authored independent fixtures, controlled experiments for replay drift, leakage, and counterfactual augmentation, plus two deterministic USS pilots outside robotics. The empirical claims in this release are intentionally bounded: they demonstrate executable failure mechanisms and auditable conversion semantics, not a completed multi-lab benchmark, production game engine audit, autonomous-driving fleet audit, or real-robot deployment study.

2 Problem Statement

USS starts from a practical question: what must be known before an interactive spatial state can be safely converted, replayed, synchronized, augmented, or used for evaluation? A valid state record must answer five questions. A production dataset must also answer how those records are named, sharded, indexed, resolved, and versioned without turning storage layout into the semantic API. WorldEpisode specializes these requirements for robot-learning episodes.

Which state revision am I in? Each record references one immutable, content-addressed state or world revision plus an ordered list of state-specific deltas. In robotics this is a base world plus episode deltas; in a game engine it may be a map patch plus live actor state; in an AV replay it may be a map/sensor-calibration revision plus log-time deltas. Changing geometry, metric scale, calibration, collision proxies, physical parameters, entity decomposition, object identity, or task-relevant fixtures creates a new revision. This makes it detectable when two datasets silently use the same filename for different states, or different filenames for the same state.

What entities does the state contain? Objects, robot links, vehicles, sensors, avatars, fixtures, support surfaces, semantic regions, and generated assets must have persistent identifiers. A mug in a segmentation mask, a Gaussian group, a collision mesh, a simulator actor, a language annotation, and a contact event should resolve to the same entity when they refer to the same physical object. A game obstacle or AV tracked object needs the same identity discipline across telemetry, assets, and runtime actors.

Where and when is this state valid? Spatial values declare their frames, units, transform direction, validity interval, and calibration revision. Streams declare clock domains, timestamps, synchronization mappings, and timing error where available. Actions declare command time and effective motor time rather than assuming those are the same instant. USS calls these state transition invariants: a control input, physics step, server patch, or sensor fusion event is only meaningful under its declared timing, frame, unit, and latency contract.

What was preserved or lost during conversion? WorldEpisode reuses existing standards rather than replacing them. USD [1] carries composed simulation worlds, glTF carries Gaussian appearance assets [14], NCore carries capture and pose-graph conventions [10], Rerun carries physical-data recordings with column chunks [12], ROS/MCAP carries logs, and domain-native engines carry their own telemetry. A converter may externalize or approximate fields, but it must emit a machine-readable loss report instead of silently discarding semantics.

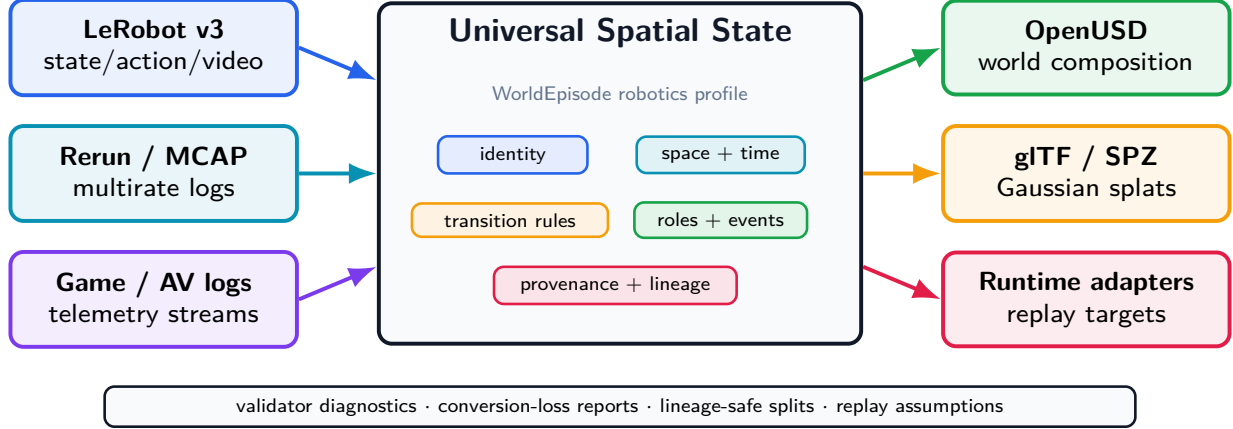


Figure 1: USS is the semantic bridge between spatial data containers, state/world assets, and runtime targets. WorldEpisode is the robotics profile used in the experiments.

Can the evaluation split leak? Record-level random splits are not enough when multiple records share a reconstructed room, map patch, physical object, source capture, calibration, generated asset lineage, or state ancestry. A USS split manifest can enforce disjointness by state lineage, physical entity, source capture, reconstruction run, generated asset family, or collection site.

Can this scale past one dataset folder? Large robot corpora require a dataset manifest above the episode records. The manifest declares namespaces, resolver priority, shard catalogs, materialized indexes, split catalogs, and append-only versions. This lets a reader locate all SO-101 pick-place traces for a world-disjoint split by index lookup rather than by crawling storage paths.

The initial WorldEpisode profile is intentionally narrow: rigid tabletop manipulation with fixed-base single- or dual-arm robots. That bounded domain keeps conformance executable rather than aspirational. USS can define later profiles for humanoids, locomotion, deformables, fluids, multi-agent game worlds, and autonomous-driving logs.

3 Universal Spatial State Architecture

USS is implemented as a sidecar blueprint around existing formats. It does not ask a LeRobot dataset to become a USD scene, a game-engine telemetry log to become an AV replay format, or a USD scene to become a policy-training table. Instead, it records the relationships those containers need in order to agree about the same state: the state revision, persistent entities, space/time conventions, transition contract, asset roles, interaction events, provenance, and split lineage. WorldEpisode is the robotics profile of this architecture.

The sidecar is not a mandate to package a global robot corpus as millions of small folders. The single-episode record is the semantic unit. Production datasets add a manifest layer: globally scoped namespaces, URI resolvers, shard catalogs, materialized indexes, split manifests, and append-only release snapshots. Opening such a corpus starts from the manifest and its indexes; it does not require recursively scanning an object store or loading every episode.

Figure 1 shows the intended data flow. Dataset, log, and telemetry containers plug in on the left; world assets and replay targets plug in on the right; the USS sidecar holds the semantics that would otherwise be scattered across undocumented loader assumptions. The current implementation names that robotics sidecar WorldEpisode.

This architecture gives the paper its central rule: keep native containers for the data they are good at storing, but make the cross-container assumptions explicit, checkable, and loss-aware.

3.1 USS as a Meta-Simulator Contract

USS is not tied to a single rendering or physics engine. It operates as a meta-simulator contract: a runtime becomes a trusted target by implementing an adapter that preserves the same semantic invariants and emits replay evidence. This keeps the integration direction clean. Datasets are not rewritten around simulator-native assumptions; simulators compile their scene loading, action interfaces, and replay reports against the USS sidecar.

The adapter contract has three layers. First, the *invariant interface* ingests immutable world revisions, frame and clock graphs, entity identity, representation roles, action channels, asset digests, provenance, and quality records. Second, *asynchronous schema extension* allows an engine to register cloth, deformable, fluid, articulation, material, or procedural-scene extensions without weakening the core episode record. Third, *deterministic replay accountability* records runtime identity, version, solver, timestep, actuator parameters, initialization state, command/effective timing, latency, tolerance envelopes, and measured drift.

The current implementation reflects this boundary across two codebases. WorldEpisode has tested minimal MuJoCo and Genesis replay adapters for the same LeRobot control-loop trace. URDF Studio, the companion simulator-transfer workbench, implements the same episode-backend idea for MuJoCo and Genesis: both backends pass a shared **SimBackend** conformance suite, and a one-episode carton-sorting scenario produces a MuJoCo–Genesis comparison report with matching task success and measured trajectory divergence. The Isaac mapping is emitted and deployment-guarded but untested, and SAPIEN remains guarded rather than replay-tested. The claim is therefore not simulator-independent physics. The claim is runtime-neutral accountability: every simulator reports against the same state contract, so cross-runtime drift becomes measurable rather than hidden inside loader conventions.

4 Five Questions, Five Graphs

Figure 2 shows the five coupled graphs that make a spatial state replayable and auditable rather than just stored. Each graph answers a question that appears during debugging: what object is this, where and when is it, what is the asset safe to do, what physically or virtually happened, and where did this data come from?

4.1 What object am I looking at?

Every physical or logical entity receives a persistent identifier: robot, link, gripper, vehicle, avatar, sensor, rigid object, articulated object, fixture, support surface, region, human, light, or semantic group. The same `entity_id` must be used by segmentation masks, bounding boxes, Gaussian groups, render meshes, collision bodies, simulator actors, language references, contact events, object trajectories, grasp events, and attachment events. The real-world fix is simple: if a language command says “pick up the mug,” the perception mask, the simulator object, and the collision proxy must identify the same mug.




Figure 2: The five-graph model. USS makes identity, frames/clocks, state roles, temporal events, and provenance mutually referential so conversion, synchronization, and replay errors become detectable instead of hidden.

4.2 Where is it, and exactly when?

Every spatial value declares source frame, target frame, units, handedness, up axis, transform direction, quaternion convention, validity interval, calibration revision, and uncertainty where available. Every stream declares a clock domain, device timestamp, host-receive timestamp where available, exposure interval for cameras, effective command time for actions, synchronization mapping, and estimated offset/drift/error. NCore’s explicit pose-graph convention demonstrates the kind of rigor USS requires across the whole state record [10]. This graph is what turns a suspicious replay from “the policy is bad” or “the engine is wrong” into a diagnosable offset, transform direction, or unit error.

4.3 What is this asset safe to be used for?

One physical or virtual entity can have many representations: Gaussian splat appearance, textured render mesh, metric surface mesh, depth point cloud, collision proxy, mass/inertia record, semantic affordance, telemetry stream, or learned embedding. Each representation declares its role and an asset descriptor: URI, media type, SHA-256 digest, optional mirrors, optional embedded payload metadata, coordinate frame, units, validity interval, derivation, uncertainty, and license. The URI may be relative, HTTP(S), Hugging Face, object storage, OCI, IPFS, or another registered resolver. Portability comes from deterministic resolution and digest verification, not from requiring every asset to live in one folder. The role field is essential: a splat may be valid for appearance but invalid for collision, while a simplified proxy may be valid for collision but unsuitable for visual training.

4.4 Did the robot actually grab the object?

Spatial records do not store only poses. They record interaction events: contact begin/end, grasp establish/release, attachment create/remove, support, containment, map or collision patch, articulation state changes, spawn/despawn, topology changes, simulator reset, intervention, subgoal completion, and failure detection. Each event references entities, time, confidence, source, and optionally evidence. This turns raw coordinate streams into milestones such as “contact established,” “object dropped,” “authoritative collision patch applied,” or “intervention occurred.”

4.5 Where did this data come from?

Every derived asset identifies source episodes, source sensors and intervals, reconstruction algorithm, software version, model/checkpoint hash, configuration, calibration revision, manual edits, source and output hashes, creator, license, and quality metrics. USS maps to existing provenance and dataset metadata vocabularies rather than inventing an incompatible publication vocabulary. This graph lets a split avoid putting the same reconstructed room, capture session, or generated asset family on both sides of evaluation.

5 Reference Implementation

5.1 Complete Action Semantics

An action vector is not portable without its operational contract. Each action channel must declare actuator, control mode, parameterization, reference frame, units, absolute/delta/velocity semantics, limits, command rate, command timestamp, effective timestamp, latency model, interpolation, normalization, gripper semantics, and missing-value policy. A vector such as $[0.01, 0, 0, 0, 0, 0, 1]$ is otherwise ambiguous: it could be a tool-frame delta, base-frame delta, velocity command, normalized policy output, or delayed target increment.

5.2 Immutable World Revisions

Each episode references one base world revision plus ordered episode deltas. A new revision is required when world geometry, metric scale, calibration, collision proxies, physical parameters, entity decomposition, object identity, or task-relevant fixtures change. Content hashes allow two datasets to prove that they refer to the same world without relying on filenames.

5.3 Loss-Aware Conversion

WorldEpisode supports many bindings, but it does not promise impossible universal losslessness. Every converter emits a machine-readable report listing preserved, externalized, approximated, discarded, and warning fields. The standard permits lossy conversion; it forbids silent lossy conversion.

5.4 Implementation Artifact

We implement the contract as four executable components: a JSON Schema for structural validity, a machine-readable conformance requirement file, an installable semantic validator, and a round-trip experiment runner. The validator maps each diagnostic to a stable requirement identifier such as `TIME.001`, `FRAME.002`, or `ASSET.002`. The experiment runner takes a WorldEpisode package, exports binding-native artifacts plus sidecars, reimports them, compares a versioned semantic

projection, injects controlled faults, and writes reproducible result artifacts. The projection itself is published as `conformance/projections/uss-core-23.v0.json` and validated against `schemas/semantic-projection-v0.schema.json`. URDF Studio remains the initial implementation root, but the contract is represented independently in the public schema and validator.

The validator is packaged as `worldepisode`. A blocking preflight can be run from the shell with `worldepisode preflight <dataset-or-manifest>` or from a training script with a single Python call such as `preflight_lerobot(dataset.root).raise_if_failed()`. Native LeRobot folders and Rerun `.rrd` recordings are recognized directly; without a WorldEpisode manifest or sidecar, they fail closed on missing world revision, entity identity, representation role, action timing, frame/clock, split-lineage, and conversion-loss controls. This makes the semantic audit a pre-training check rather than a separate post-hoc documentation exercise. The dataset-scale audit checks the separate corpus manifest for namespace uniqueness, resolver coverage, digest-addressed assets, local mirrors, shard/index references, required lineage and asset-digest indexes, split-manifest presence, and append-only version structure. A companion scale-performance benchmark generates a 32,768-shard descriptor catalog with 1,073,741,824 described episodes and measures catalog open/indexing, partition pruning, digest-cache resolution, and resolver routing without materializing payload rows. The clean-room reader check intentionally does not import the `worldepisode` package. It uses the published JSON Schema and separately implemented requirement checks to parse the minimal example and fixture corpus, providing an internal sanity check for independent implementability.

The repository also includes an active LeRobotDataset v3 converter. It downloads bounded Parquet and metadata shards from the pinned public `lerobot/svla_so101_pickplace` SO-101 dataset, constructs a schema-valid WorldEpisode manifest and sidecar, exports a small LeRobotDataset v3 package, reads that package back, and asserts equality of source-native tensors and timestamp records. Source fields that LeRobot does not declare, such as camera extrinsics and controller latency, are not guessed; they are emitted in the conversion report as explicitly source-absent. The split-leakage tool similarly downloads bounded ArmnetBench LeRobot shards, writes `world_lineage` and split-manifest artifacts, and trains an executable BC baseline over the real LeRobot tensors. The control-replay tool reads the exported LeRobot trajectory, estimates the effective command delay from action/state alignment, writes the WorldEpisode action contract, executes tested MuJoCo and Genesis adapters, and emits an Isaac adapter mapping that is ready for an Isaac environment but explicitly marked untested in the committed report. The replay-adapter conformance tool then runs a dependency-free reference scheduler over delay, zero-order-hold, missing-command, and asynchronous queue cases. This checks the timing contract that runtime adapters must honor before a simulator-specific replay is trusted.

The repository also emits a meta-simulator adapter report. It defines three compliance layers: the invariant interface, asynchronous schema extension, and deterministic replay accountability. The report records WorldEpisode’s tested MuJoCo and Genesis replay adapters, URDF Studio’s tested MuJoCo/Genesis episode-backend conformance and one-episode comparison, Isaac as adapter-ready but untested, and SAPIEN as adapter-required. This artifact is deliberately an adapter contract rather than a claim of equal runtime coverage.

6 Failure Case Studies

The experiments are written as failure investigations rather than as isolated unit tests. Each one starts from a failure mode that can make an embodied or virtual-agent result look better, cleaner, or more reproducible than it really is. Robotics provides the deep measured stress test; two deterministic USS pilots check that the same state-invariant vocabulary also applies outside robot

episodes.

How to read amber boxes. Amber boxes mark work that is executable but not counted as a paper result. Each box states the exact command path, required artifacts, and acceptance rule needed before the corresponding claim can move from “open” to “measured.” These boxes are intentionally visible in the PDF so a reader can reproduce or close the missing evidence without reading the repository internals.

Case 0: the valid file with the wrong live state. USS is evaluated outside robotics with two deliberately small pilots. In the game-engine collision case, a client can still load an old arena collision asset, but an authoritative patch moves the obstacle into the avatar path. The stale client path looks valid locally, fails against the authoritative state, and is caught by the asset digest, state revision, and role checks (`ASSET.002`, `WORLD.001`, `REP.001`). In the autonomous-driving clock case, camera and lidar logs deserialize correctly, but a 50 ms undeclared clock offset at 15 m/s creates a 0.75 m fusion error against a 0.20 m tolerance. The USS clock mapping corrects the error to 0.0 m and triggers `TIME.001`, `TIME.002`, and `FRAME.001`. These pilots support the terminology shift from episodes to state; they are not production game-engine or AV benchmark results.

Case 1: the split that faked success. The most dangerous failure is not a crash; it is a benchmark that looks successful for the wrong reason. We audit ArmnetBench v0.1’s single-arm SO-101 LeRobot release [2] by deriving eight WorldEpisode-style `world_lineage` hashes, one for each task-scene and camera-layout group. A standard random episode split trains on 320 episodes, tests on 80, and leaks all test scene lineages into training. The same Torch MLP state-action behavioral-cloning probe then reaches 0.850 offline imitation success under a fixed normalized-RMSE tolerance of 0.25.

When the split is made scene-disjoint, the apparent success disappears. The held-out eye-drop scene lineages have zero overlap with training, the test set contains 100 episodes, and offline BC success drops to 0.000. Mean episode normalized RMSE rises from 0.183 to 0.363, a $1.98\times$ increase. This is an offline action-imitation benchmark, not a physical rollout claim or a claim about state-of-the-art LeRobot policies. Its purpose is narrower and diagnostic: it exposes that a conventional split can reward memorizing a background world rather than learning a transferable manipulation behavior. The committed split manifest makes the same leakage audit repeatable with ACT, Diffusion Policy, or another LeRobot-native policy.

To verify that the split packages are executable beyond the original MLP probe, we also run a closed-form temporal ridge baseline over the committed compact LeRobot state/action packages. The baseline uses a three-frame state history and no video. It reaches 0.925 offline success on the random episode split and 0.420 on the scene-disjoint split, a 0.505 success-rate drop, while mean episode normalized RMSE rises from 0.157 to 0.255. This is still not ACT or Diffusion and not a rollout result, but it shows that the leakage penalty survives a second temporal policy-style baseline executed directly on the packaged LeRobot splits.

To keep that stronger claim from being implied before it is measured, the repository also emits an ACT/Diffusion gate from the same split manifest. The gate writes train/test episode allowlists, four virtual split manifests with 27 digest-verified source files, four compact physical LeRobot state/action split packages totaling 241,470 rows, four LeRobot-native training jobs (`act` and `diffusion` on random and scene-disjoint splits), required offline action-evaluation reports, and a high-fidelity simulator or physical rollout contract. The gate is intentionally open in this release; no

ACT, Diffusion Policy, IsaacLab, or hardware success number is claimed, and the compact packages omit video payloads until source videos are mirrored with committed digests.

Open result, not claimed: ACT/Diffusion and rollout impact

Status: open. To close this gate, first regenerate the split packages and LeRobot job specs with `python3 tools/lerobot_policy_leakage_gate.py`. Then run the generated jobs with `bash docs/experiments/lerobot_policy_gate/run_lerobot_policy_jobs.sh` inside a LeRobot training environment.

Required artifacts: policy checkpoint digests, train/eval configs and seeds, offline action metrics for random and scene-disjoint splits, rollout traces or videos with content digests, and an updated `docs/experiments/lerobot_policy_gate/policy_gate_report.json`.

Acceptance rule: at least one ACT or Diffusion Policy run must report both split protocols, and at least one simulator or physical rollout report must use the same split manifest. Until then, the paper claims only the offline Torch probe and the prepared gate.

Figure 3 summarizes the measured changes for three experiment families where explicit state semantics alter the outcome.




Figure 3: Quantitative summary of the three measured effects used in the evaluation. Lower is better for leakage and replay RMSE; higher is better for semantic retention.

Case 2: the converter that refuses to invent missing physics. To test an actual LeRobot-Dataset rather than only a capability model, we implement LeRobotDataset v3 $\rightarrow$ WorldEpisode $\rightarrow$ LeRobotDataset v3 on five-episode batches from two pinned public datasets: `lerobot/svla_so101_pickplace` [8] and `lerobot/pusht` [7]. The converter downloads the real Parquet metadata and data shards, builds a schema-valid WorldEpisode manifest and sidecar, exports a compact LeRobotDataset v3 package, reads it back, and asserts exact equality.

This is a two-dataset batch audit for active conversion, not a coverage claim for all LeRobot datasets. Across ten episodes, the round trip preserves 1,935 rows of source-native action tensors, 1,935 state rows, sample timestamps, frame indices, episode indices, global sample indices, task indices, and camera video timestamp ranges with maximum absolute error 0.0. It also preserves the derived physical-frame records through the sidecar. More importantly, it does not fake missing semantics: camera extrinsics, robot/world calibration, action units, and controller latency are reported as source-absent rather than guessed. This is the difference between a convenient conversion and an auditable one.

Case 3: the replay that drifted because time was underspecified. LeRobot’s control stack can decouple policy inference from the fixed-rate low-level controller: commands may be

queued, chunked, interpolated, or consumed after policy output [3]. On the public SO-101 trajectory, WorldEpisode records policy-output, enqueue, queue-consume, and motor-receive timestamp semantics, plus zero-order-hold interpolation and a constant-frame latency model.

Estimating the effective delay on the calibration prefix gives four 30 Hz frames, or 133 ms. On the validation suffix, naive command-time alignment has joint RMSE 4.732 degrees. Timestamp-aware alignment has RMSE 1.862 degrees, a $2.54\times$ improvement. We then execute the same WorldEpisode action contract through two minimal six-joint position-servo replay adapters. In MuJoCo, naive command-time replay has RMSE 3.425 degrees and timestamp-aware replay has 1.563 degrees, a $2.19\times$ improvement. Genesis same-trace replay gives the same measured profile: 3.425 degrees under naive command-time replay and 1.563 degrees under timestamp-aware replay, also a $2.19\times$ improvement. The Isaac adapter is emitted as a ready mapping from the same contract, but is not tested in this environment. The failure avoided here is a common one: blaming a simulator or controller when the dataset never said when the command actually took effect.

We also add a dependency-free replay-adapter conformance harness before trusting additional runtime adapters. It tests three scheduler cases: constant frame delay, zero-order hold for missing commands, and asynchronous queue selection. In all three cases a naive scheduler fails while the contract-aware scheduler matches the expected trace with zero numerical error. This is a precondition check for simulator adapters, not a claim that the simulators have equivalent physics. For broader runtime evidence, we also record URDF Studio’s shared **SimBackend** conformance suite for MuJoCo and Genesis and its one-episode carton-sorting MuJoCo–Genesis comparison report with matched task success and measured trajectory divergence. These results close the narrow same-trace second-runtime replay gate for the minimal joint replay profile. They do not claim contact-rich task replay, Isaac/SAPIEN execution, or simulator-independent behavior.

Case 4: real-to-sim contract drift. Recent Gaussian/OpenUSD real-to-sim systems correctly separate visual reconstruction from interactive physical assets. WorldEpisode targets the contract between those layers and the source robot episode. We therefore add two deterministic ablations. In the action-contract ablation, a policy succeeds in a simulator that interprets its output as absolute joint-radian targets, but the same vector fails under a deployment proxy whose controller expects delta-degree commands. In the representation-role ablation, an appearance-only Gaussian export drops the collision proxy; the straight-line policy succeeds in simulation and collides with the real foreground geometry under the deployment proxy. In both ablations the drifted contract has 2/2 simulated successes and 0/2 deployment successes, while the WorldEpisode contract recovers 2/2 deployment successes. This is a controlled proxy, not a RoboSnap or hardware rerun, but it shows why visual reconstruction is not enough without action contracts and representation roles.

Case 5: the validator that catches physically impossible packages. We inject 14 faults into the baseline package: missing clock domains, incomplete cross-clock mappings, unknown frames, missing transform intervals, unknown event entities, missing representation roles, digest errors, action-contract omissions, non-content-addressed worlds, missing delta sequences, missing provenance, and missing quality records. The validator detects all 14 expected requirement failures with 1.000 recall and 0.933 precision. Two hand-authored independent fixtures also trigger the expected diagnostics. The single false-positive requirement is useful rather than harmful: removing an action reference frame triggers both **ACTION.001** and the downstream **FRAME.001** reference check.

The same suite includes a simple unit sanity check from `lerobot/svla_so101_pickplace`. The first mirrored SO-101 frame has six joint values with maximum magnitude 98.924 in degrees and 1.727 after radian conversion. Treating those degree values as radians would violate the action/state

unit contract even though the tensor shape is valid.

To move beyond synthetic mutations, the experiment runner also writes a pilot natural-source corpus from committed public LeRobot artifacts and selected source-level public benchmark metadata. Across `lerobot/svla_so101_pickplace`, `lerobot/pusht`, `armnet/armnetbench_v01_`, `lerobot_so101`, DROID, and BridgeData V2, it records 19 observed cases. The active LeRobot evidence covers source-absent camera extrinsics, robot/world calibration transforms, action units, controller-latency models, and a random episode split with measured scene-lineage leakage. The DROID and BridgeData V2 cases are source-level public evidence gaps selected from the famous-benchmark call-out audit. These are not claimed as maintainer-confirmed dataset bugs or measured score-inflation results; they are naturally occurring source omissions, metadata gaps, or evaluation risks that a WorldEpisode audit must make explicit.

The same validator is exposed as a pre-training command rather than only as an experiment script. The preflight regression covers four cases: a valid WorldEpisode manifest passes, an invalid WorldEpisode fixture fails, a native LeRobot v3 folder without a WorldEpisode sidecar fails closed, and a Rerun `.rrd` recording without a sidecar fails closed. This is the intended usage pattern for labs: run one blocking line before launching a multi-day policy training job. To reduce single-implementation risk, we also run an internal clean-room reader that does not import the `worldepisode` package. It parses the public schema, summarizes the minimal example, and catches all 18 expected requirement instances across 17 pilot and independently authored fixtures with 1.000 recall. This is not external adoption, but it verifies that the public artifacts are readable outside the reference SDK path.

Case 6: the corpus that should not require a filesystem crawl. To address production scale, we audit the dataset manifest in `examples/scalable-corpus.worldepisode-dataset.json`. The audit validates five namespaces, three resolver schemes, three registries, three shards, two indexes, one append-only version, and nine digest-addressed asset descriptors spanning Hugging Face, S3, and OCI URIs with local mirrors. It also checks that the manifest includes a world-lineage index, an asset-digest index, and a split-manifest shard. We then run a generated catalog benchmark with 32,768 trace shards describing 1,073,741,824 episodes. The committed report records the measured time for parsing and indexing the generated JSON catalog, eight partition-pruning queries, digest-cache resolution, and resolver routing over all asset descriptors. This is catalog-side evidence only, but it makes the scale claim executable: opening a corpus starts from the manifest and indexes rather than from recursive storage discovery.

Case 7: the famous benchmarks that still need call-out audits. The ArmnetBench result is the measured leakage case. To scale the same diagnostic, we add a source-level call-out audit over five famous public robot-learning benchmarks: Open X-Embodiment/RT-X [11], DROID [5], BridgeData V2 [16], LIBERO [9], and CALVIN [4]. The audit asks whether public metadata exposes world or scene lineage splits, content-addressed world revisions, persistent entity identity, action timing and latency contracts, frame/clock graphs, and replay or conversion-loss reports. All five benchmarks have at least one high-severity open control; the three real-robot datasets have five each in the current source-level audit. This is a call-out target list, not a claim that their published scores are inflated. The stronger claim requires converting each benchmark into WorldEpisode and rerunning a published policy protocol under lineage-disjoint splits or timestamp-aware replay. To make that boundary executable, we add a separate inflation-proof gate: it requires a benchmark-specific WorldEpisode conversion, lineage/timing audit, published-policy rerun, corrected evaluation, and measured score drop before any famous benchmark can be labeled inflated. We also add a

targeted DROID-100 LeRobot-subset rerun tool that downloads pinned public Parquet shards, builds a WorldEpisode evidence package, compares random and lineage-disjoint splits, and runs a deterministic offline state-action policy. The gate records one attempted DROID subset rerun report, zero valid famous-benchmark rerun reports, and zero measured famous-benchmark inflation claims in this release.

Open result, not claimed: famous-benchmark score-inflation proof

Status: open. A targeted DROID subset rerun can be attempted with `uv run -with pyarrow -with requests -with numpy python tools/famous_benchmark_policy_rerun.py -benchmark droid_100 -required`. The proof contract is then enforced with `python3 tools/benchmark_inflation_gate.py -required`.

Required artifacts: benchmark-specific WorldEpisode conversion, lineage/timing audit, faithful published-protocol policy rerun, paired corrected evaluation under the same metric, and a measured baseline-minus-corrected score drop.

Acceptance rule: the gate must contain at least one valid rerun report with `measured_inflation=true`. Source-level metadata gaps alone are evidence of missing controls, not evidence that published scores are inflated.

Case 8: augmentation that uses the world rather than a side file. We evaluate a shifted-camera/object-displacement regression task. Observations-only nearest-neighbor control reaches the proxy success threshold on 0.292 of shifted test cases. Adding unstructured 3D side information increases proxy success to 0.342. WorldEpisode-style world-frame identity and counterfactual camera augmentation increase proxy success to 0.592 over 120 deterministic test cases.

Binding retention. The same executable artifacts measure how much of the semantic contract is preserved by native containers alone. Each binding writes a native payload and a WorldEpisode sidecar, then the experiment reimports both and checks exact equality against the versioned `uss-core-23` projection in `conformance/projections/uss-core-23.v0.json`: episode identity/task outcome, world revision and asset descriptor, embodiment identity and URDF asset, frames/transforms, clocks/mappings, entity identity/roles/assets, action control and timing contracts, trace binding and asset descriptor, events, ordered deltas, provenance, quality, split lineage, and replay assumptions. This is an executable pilot projection, not a universal score of each format’s value. Across seven pilot bindings, Table 1 shows that common containers preserve 17–39% of this projection natively, while WorldEpisode sidecars recover that projection for the tested dataset/log/world bindings. A Gaussian asset binding alone preserves 17% natively and 30% with limited sidecar metadata, confirming that an asset format is not an episode–world contract.

Table 1: Semantic field retention in the pilot binding model.

Binding	Native	With WorldEpisode
LeRobot v3	0.261	1.000
Rerun RRD	0.348	1.000
NCore	0.304	1.000
MCAP/ROS2	0.261	1.000
OpenUSD/SimReady	0.391	1.000
glTF Gaussian asset	0.174	0.304

Finally, we generate episodes with world/entity lineage labels and evaluate a memorization

policy. A random episode split leaks every tested world/entity pair and obtains 1.000 accuracy. World-disjoint and entity-disjoint splits have zero measured leakage and reduce memorization accuracy to 0.500 and 0.562 respectively. This pilot demonstrates why split manifests need lineage constraints beyond the public ArmnetBench audit above.

7 Limitations and Scope

The current evidence is deliberately scoped. It is meant to expose hidden failure mechanisms, not to claim full robot-learning, game-engine, or autonomous-driving generality.

Table 2 states the boundary of each empirical claim. This is important because the paper argues for a semantic contract, not for a finished benchmark suite.

Table 2: What the current evidence supports, and what it does not yet support.

Claim area	Current evidence	Boundary
Leakage	Public ArmnetBench LeRobot audit with 400 teleoperated reference episodes, an executable Torch BC probe, a measured temporal ridge state/action baseline, an ACT/Diffusion gate harness, and four compact digest-verified LeRobot state/action split packages totaling 241,470 frames	Stronger ACT/Diffusion jobs are prepared but not run; the compact split packages omit video payloads, so no vision-policy, real-robot, or high-fidelity rollout result is claimed yet
Conversion	Two pinned public LeRobotDataset v3 five-episode batch round trips with exact tensor, index, and timestamp equality	Two datasets; broader LeRobot coverage remains future work
Replay timing	Real SO-101 trajectory alignment, tested same-trace MuJoCo and Genesis position-servo replay adapters, URDF Studio MuJoCo/Genesis episode-backend evidence, and a dependency-free scheduler conformance harness	One LeRobot trace with minimal joint replay adapters; no contact-rich task rollout, Isaac runtime result, or SAPIEN runtime result is claimed
Meta-simulator neutrality	Adapter contract matrix over MuJoCo, Isaac Sim, Genesis, and SAPIEN plus same-trace MuJoCo/Genesis replay evidence, URDF Studio SimBackend conformance, and a one-episode MuJoCo-Genesis comparison	MuJoCo and Genesis are tested for the minimal replay profile; Isaac and SAPIEN are not replay-tested in this paper, and equal physics is not claimed
USS generality	Deterministic game-engine collision-patch and autonomous-driving clock-domain pilots using the same invariant vocabulary	Not measured on production game engines, public AV datasets, or fleet logs
Validation	Fourteen injected requirement faults, two independent hand-authored fixtures, and a pilot natural-source corpus over five public datasets	Five-dataset count is met through active LeRobot artifacts plus source-level public benchmark metadata; no maintainer feedback yet
Binding retention	Versioned uss-core-23 semantic projection checked by executable artifacts	Pilot projection; not a universal score of each storage format
Famous benchmark call-out	Source-level audit over Open X-Embodiment, DROID, BridgeData V2, LIBERO, and CALVIN, plus a targeted DROID subset rerun tool and a separate inflation-proof gate	Missing public controls are not measured score inflation; the proof gate currently has one attempted DROID subset rerun report, zero valid famous-benchmark rerun reports, and zero measured inflation claims
Real-to-sim drift	Controlled action-contract and representation-role ablations where drifted contracts succeed in simulation and fail under deployment proxies	Deterministic proxy; not a RoboSnap/DROID-Sim rerun or physical hardware deployment
Adoption	Public schema, installable local preflight package, validator, fixtures, governance files, and an internal clean-room reader that does not import the reference SDK	No PyPI release, upstream LeRobot/Rerun integration, external independent implementation, or external dataset release yet
Scale architecture	Dataset manifest schema, scalable example, executable audit, and a generated 32,768-shard catalog benchmark describing 1,073,741,824 episodes for open/index latency, partition pruning, digest-cache behavior, and resolver routing	Catalog-side benchmark only; no billion episode rows, payload bytes, distributed storage, concurrent readers, or institutional federation are measured

Open result, not claimed: results not claimed in this release

The artifact keeps unfinished claims executable instead of hiding them. The current RFC release gate is: `python3 tools/release_readiness.py -strict-rfc`. The full-standard gate is: `python3 tools/release_readiness.py -strict-full`. The latter is expected to fail until all open gates below are backed by committed reports. The same gates are indexed in `docs/experiments/open_reproduction_gates/open_reproduction_gates.json`.

- ACT/Diffusion leakage: run `python3 tools/lerobot_policy_leakage_gate.py`, then execute `docs/experiments/lerobot_policy_gate/run_lerobot_policy_jobs.sh`. A paper update must commit the policy metrics, rollout traces or videos, and updated `docs/experiments/lerobot_policy_gate/policy_gate_report.json`.
- Famous-benchmark inflation: run the targeted DROID subset rerun with optional Arrow/HTTP dependencies, using `tools/famous_benchmark_policy_rerun.py` with `-benchmark droid_100 -required`, then require `python3 tools/benchmark_inflation_gate.py -required`. No benchmark may be described as inflated unless this gate records a valid rerun and a measured corrected-score drop.
- External standard evidence: add maintainer-reviewed diagnostics, an independently written reader/exporter, and at least one externally published USS or WorldEpisode-compatible dataset before calling the project a mature interoperability standard.

Manipulation scope. The current profile targets rigid tabletop manipulation with fixed-base single- or dual-arm robots. Humanoids, locomotion, deformables, fluids, and multi-agent environments require later profiles and new conformance fixtures.

USS generality. The paper title and framing are broader than robotics, but the strong empirical evidence is still the WorldEpisode robotics profile. The game-engine and autonomous-driving examples are deterministic pilots that test whether the same state-invariant vocabulary can express non-robotics drift. They do not prove prevalence or effectiveness on production game, simulation, or AV datasets. A stronger USS paper should add at least one public non-robotics corpus and an independently maintained adapter.

Offline leakage result. The ArmnetBench experiment is an offline behavioral-cloning audit, not a physical rollout. Its value is that it demonstrates how random episode splits can leak world lineage and inflate evaluation, not that it measures real-robot success directly. The current baseline is an executable Torch MLP probe over LeRobot tensors; ACT and Diffusion Policy should be run on the committed split manifest before claiming impact on common LeRobot training recipes. The repository now also measures a temporal ridge baseline over the committed compact packages; it drops from 0.925 random episode offline success to 0.420 scene-disjoint success, but remains a low-dimensional offline state/action result. The repository emits ACT/Diffusion jobs, compact state/action LeRobot split packages, and rollout requirements, but no ACT/Diffusion policy result is claimed here. The split packages verify 27 cached source files, preserve state/action/timestamp rows, and remap episode/index fields into contiguous local LeRobot folders; they intentionally omit video payloads, so a vision-policy result still requires mirrored source videos with committed digests.

Replay is not bit-identical physics. WorldEpisode records simulator assumptions, timing, interpolation, and tolerances; it does not promise simulator-independent behavior. The reported replay result is a same-trace MuJoCo/Genesis position-servo test under declared assumptions. The replay-adaptor harness checks delay, zero-order hold, missing-command, and asynchronous queue scheduling, but it is not a physics simulator. Isaac and SAPIEN are intentionally untested in this environment, and no contact-rich task rollout result is claimed.

Meta-simulator evidence. The meta-simulator section defines what a compliant runtime adapter must preserve. It does not certify that MuJoCo, Isaac, Genesis, SAPIEN, or any other simulator has equivalent physics. Current evidence contains WorldEpisode MuJoCo and Genesis replay adapters for the same LeRobot control trace, plus a separate URDF Studio episode-backend path where MuJoCo and Genesis pass shared conformance tests and produce a one-episode comparison report. Isaac and SAPIEN are not replay-tested here.

Reference implementation maturity. The converter, validator, and reports are executable, but independent implementations are still needed before the project should be treated as a mature interoperability standard. The clean-room reader reduces the risk that the fixtures only work through the reference SDK, but it is still an internal artifact. The next empirical step is to run the same protocol on larger multi-camera manipulation datasets and externally maintained adapters.

Preflight distribution. The repository packages a local `worldepisode` preflight CLI and Python API, but it is not yet a released PyPI package and has not been merged into upstream LeRobot or Rerun examples. The current claim is friction reduction in the reference implementation, not ecosystem default status.

Dataset-scale performance. The specification separates the core episode record from a dataset-scale manifest with namespaces, shards, indexes, resolvers, and append-only snapshots. That design prevents the standard from being defined around a local folder layout. The current audit makes the catalog invariants executable, and the scale-performance benchmark opens and indexes a generated 32,768-shard descriptor catalog with 1,073,741,824 described episodes, then measures partition pruning, digest-cache resolution, and resolver routing. This is still not a measured distributed-systems result: it does not materialize the episode rows, load payload bytes, test concurrent readers, measure cache eviction, or exercise federation across institutions.

Famous benchmark call-out scope. The audit over Open X-Embodiment, DROID, BridgeData V2, LIBERO, and CALVIN is a source-level metadata audit. It identifies where public materials do not expose the controls needed to test world-lineage leakage, action latency, or replay loss. It does not prove any published benchmark score is inflated. That claim requires a benchmark-specific WorldEpisode conversion and a measured policy rerun. The repository now makes that requirement executable with a targeted DROID subset rerun tool and a separate proof gate. In the current artifact, one DROID subset rerun was attempted, zero valid famous-benchmark rerun reports are committed, and zero score-inflation claims are measured.

Real-to-sim drift scope. The contract-drift experiment isolates two real-to-sim failure mechanisms: action-interface drift and missing representation roles. It is intentionally deterministic and does not use a reconstructed RoboSnap/DROID-Sim scene or a physical robot. The stronger result is to bind a public reconstructed scene, replay the same checks, and measure simulator/deployment divergence under declared tolerances.

Natural failure corpus. The injected conformance faults prove that requirements are executable and diagnosable. The current pilot natural-source corpus adds 19 cases from five public robot-learning datasets, but it still does not prove prevalence in the wild. Two datasets are active LeRobot conversion audits, one is an active LeRobot scene-leakage audit, and two are source-level public metadata audits from the benchmark call-out list. A stronger version of this paper should convert

those source-level gaps into dataset-specific WorldEpisode manifests, report false-positive review, and record whether dataset maintainers agree with representative diagnostics.

8 Discussion

Why this is not only infrastructure. The experiments show several ways a result can be wrong while the files are valid: a split can leak world lineage, replay can drift because control timing is underspecified, conversion can discard physical semantics that the target container has no native place to store, and a live virtual or AV state can drift from its authoritative revision. USS is useful only because these failures are observable and testable. The sidecar is the mechanism; the scientific claim is that explicit spatial-state semantics change what an embodied-agent result means. The contribution is therefore a set of falsifiable invariants: persistent entity identity must be consistent across observations and assets, timestamps must distinguish observation, command, and effective time, state roles must separate appearance from physics, and evaluation splits must be able to exclude shared state lineage. A container adapter either preserves those invariants, externalizes them with declared loss, or fails conformance.

Not another monolithic binary. USS has a small semantic core and multiple bindings. WorldEpisode is the robotics profile: LeRobotDataset v3 is the policy-training view; Rerun is the multirate physical-data view; NCore is the sensor/reconstruction capture view; MCAP/ROS is the raw log view; OpenUSD/SimReady is the composed simulation-world view; glTF or SPZ is the Gaussian asset view; and a reference directory is the conformance profile. OpenUSD standardizes how the 3D world is composed, but USS standardizes how any agent, whether a physical robot, a video game character, or an autonomous vehicle, modifies state within that space over time without silent data corruption. A game-engine or AV profile should reuse the same invariant structure rather than inherit robotics-specific terminology.

Real-to-sim as the proving ground. Gaussian-splat and OpenUSD real-to-sim pipelines are not competitors to USS. They are the systems that most need a state and episode contract. A reconstructed room can be excellent visual training data and still be unsafe for policy evaluation if foreground collision roles are missing, if the same world lineage leaks across splits, or if the replayed policy output is interpreted under the wrong hardware action contract. USS, through WorldEpisode, should be judged by whether it keeps those pipelines auditable as they scale, not by whether it serializes splats better than asset formats designed for that job.

Simulators as accountable adapters. The same principle applies to physics and rendering engines. USS should not become a MuJoCo wrapper, an Isaac wrapper, or a Genesis wrapper. It should define the contract those adapters must satisfy: ingest the invariant interface, declare extensions, and report replay assumptions and drift. This is a stronger runtime-neutral stance than claiming universal behavior. It lets a simulator team plug in by implementing the adapter, while the validator remains the common judge of whether the exported dataset is physically and temporally auditable.

Lineage-safe evaluation. Record-level random splits can leak the same reconstructed room, physical object, Gaussian asset, source video, state revision, map patch, or generated asset lineage across training and evaluation. USS supports split constraints such as `world_lineage`, `physical_entity`, `source_capture`, `reconstruction_run`, and `generated_asset_family`.

Call-outs without overclaiming. The practical target is to make famous benchmarks auditable, not to accuse them without a rerun. When public metadata cannot answer whether train and test share a scene lineage, asset family, source capture, action timing convention, or simulator assumption, WorldEpisode should mark the published score as unaudited with respect to that control. It should mark a score as inflated only after the same policy protocol is rerun under the corrected split or timing contract.

Standards alignment. The project aligns with standards bodies rather than antagonizing them. Static 3D robot-map exchange, humanoid dataset lifecycle, OpenLABEL annotation, glTF Gaussian splats, USD composition, and NCore capture conventions are all useful neighbors. WorldEpisode is an executable profile and bridge across those layers; USS is the broader state-contract vocabulary.

9 Conclusion

Spatial-agent failures often arrive disguised as clean data: valid tensors, valid videos, valid scene files, valid telemetry logs, and invalid state assumptions. USS makes those assumptions explicit. It binds interaction data to immutable state revisions, gives persistent identity to entities, declares space/time/transition semantics, assigns roles to assets, records events and provenance, and reports conversion loss instead of hiding it. WorldEpisode is the robotics profile that makes this claim executable today.

In the reference experiments, those requirements expose an 85% offline random-split result that collapses to 0% under scene-disjoint evaluation, reduce LeRobot replay drift on a real SO-101 trajectory by modeling a four-frame action delay, and catch controlled physically incoherent packages before they are used as trusted artifacts. They also show how a real-to-sim pipeline can report simulated success while failing under deployment proxies when action contracts or representation roles drift, and how the same state-invariant vocabulary catches a game-engine collision-patch drift and an autonomous-driving clock-domain drift in deterministic pilots. The validator is exposed as a one-line preflight command for WorldEpisode, LeRobot, and Rerun artifacts, and the simulator pathway is defined as an adapter compliance contract rather than a dependency on one runtime.

These are scoped results, not a completed standardization claim. The call-out audit identifies Open X-Embodiment, DROID, BridgeData V2, LIBERO, and CALVIN as priority targets for this protocol, but it does not claim their scores are inflated without reruns; the inflation-proof gate currently records one attempted DROID subset rerun, zero valid famous-benchmark rerun reports, and zero measured inflation claims. The non-robotics USS pilots are vocabulary checks, not production game-engine or AV evaluations. Gaussian splats remain a high-value appearance profile, but the durable contribution is the spatial-state contract around them. The next step is to scale the same protocol to larger multi-robot datasets, stronger LeRobot-native policies, real or high-fidelity rollout evaluation, natural failure corpora, non-robotics USS adapters, and independently maintained bindings.

References

- [1] Alliance for OpenUSD. OpenUSD: Universal scene description. <https://openusd.org/>, 2026. Accessed 2026-07-13.

- [2] Armnet. ArmnetBench v0.1: LeRobot single-arm SO-101 release. https://huggingface.co/datasets/armnet/armnetbench_v01_lerobot_so101, 2026. Pinned revision 2e5e89aee0e7f081078d9d6ab3b279fc4b83ea84; accessed 2026-07-13.
- [3] Remi Cadene, Simon Aliberts, Francesco Capuano, Michel Aractingi, Adil Zouitine, Pepijn Kooijmans, Jade Choghari, Martino Russi, Caroline Pascal, Steven Palma, Mustafa Shukor, Jess Moss, Alexander Soare, Dana Aubakirova, Quentin Lhoest, Quentin Gallouedec, and Thomas Wolf. LeRobot: An open-source library for end-to-end robot learning. In *International Conference on Learning Representations*, 2026. arXiv:2602.22818.
- [4] CALVIN Authors. CALVIN: A benchmark for language-conditioned policy learning for long-horizon robot manipulation tasks. <https://calvin.cs.uni-freiburg.de/>, 2021. Accessed 2026-07-13.
- [5] DROID Dataset Authors. DROID: A large-scale in-the-wild robot manipulation dataset. <https://droid-dataset.github.io/>, 2024. Accessed 2026-07-13.
- [6] Hugging Face. LeRobotDataset v3.0. <https://huggingface.co/docs/lerobot/lerobot-dataset-v3>, 2026. Accessed 2026-07-13.
- [7] LeRobot. pusht: A public LeRobotDataset v3 dataset. <https://huggingface.co/datasets/lerobot/pusht>, 2025. Pinned revision 7628202a2180972f291ba1bc6723834921e72c19; accessed 2026-07-13.
- [8] LeRobot. svla_so101_pickplace: A public LeRobotDataset v3 dataset. https://huggingface.co/datasets/lerobot/svla_so101_pickplace, 2025. Pinned revision f641879e22172be7e8161d5e6c1503c2d2feb657; accessed 2026-07-13.
- [9] LIBERO Authors. LIBERO: Benchmarking knowledge transfer for lifelong robot learning. <https://libero-project.github.io/main.html>, 2023. Accessed 2026-07-13.
- [10] NVIDIA. NCore Specification: Data Conventions. <https://nvidia.github.io/ncore/data/conventions>, 2026. Accessed 2026-07-13.
- [11] Open X-Embodiment Collaboration. Open X-Embodiment: Robotic learning datasets and rt-x models. <https://robotics-transformer-x.github.io/>, 2023. Accessed 2026-07-13.
- [12] Rerun. A new data layer for robot learning. <https://rerun.io/blog/data-layer-for-robot-learning>, 2026. Accessed 2026-07-13.
- [13] RoboSnap Authors. RoboSnap: One-Shot Real-to-Sim Scene Generation for Generalizable Robot Learning and Evaluation. <https://arxiv.org/html/2607.06699v1>, 2026. Accessed 2026-07-13.
- [14] The Khronos Group. KHR_gaussian_splatting. https://github.com/KhronosGroup/glTF/blob/main/extensions/2.0/Khronos/KHR_gaussian_splatting/README.md, 2026. Accessed 2026-07-13.
- [15] The Khronos Group. Khronos Announces glTF Gaussian Splatting Extension. <https://www.khronos.org/news/press/glTF-gaussian-splatting-press-release>, 2026. Accessed 2026-07-13.

- [16] Homer Walke, Kevin Black, Abraham Lee, Moo Jin Kim, Max Du, Chongyi Zheng, Quan Vuong, Philippe Hansen-Estruch, Andre He, Vivek Myers, Kuan Fang, Chelsea Finn, and Sergey Levine. BridgeData V2: A dataset for robot learning at scale. In *Conference on Robot Learning*, 2023. Accessed 2026-07-13.